http://web.cs.hacettepe.edu.tr/~aykut/classes/spring2013/bil682/tomgauld.jpg

DS501: Machine learning, Part 2

Prof. Randy Paffenroth
rcpaffenroth@wpi.edu

Worcester Polytechnic Institute

WPI

Announcements

- Midterms are being graded as we speak...

Announcements

- Case Study 3 is ready and lets have a conversation...

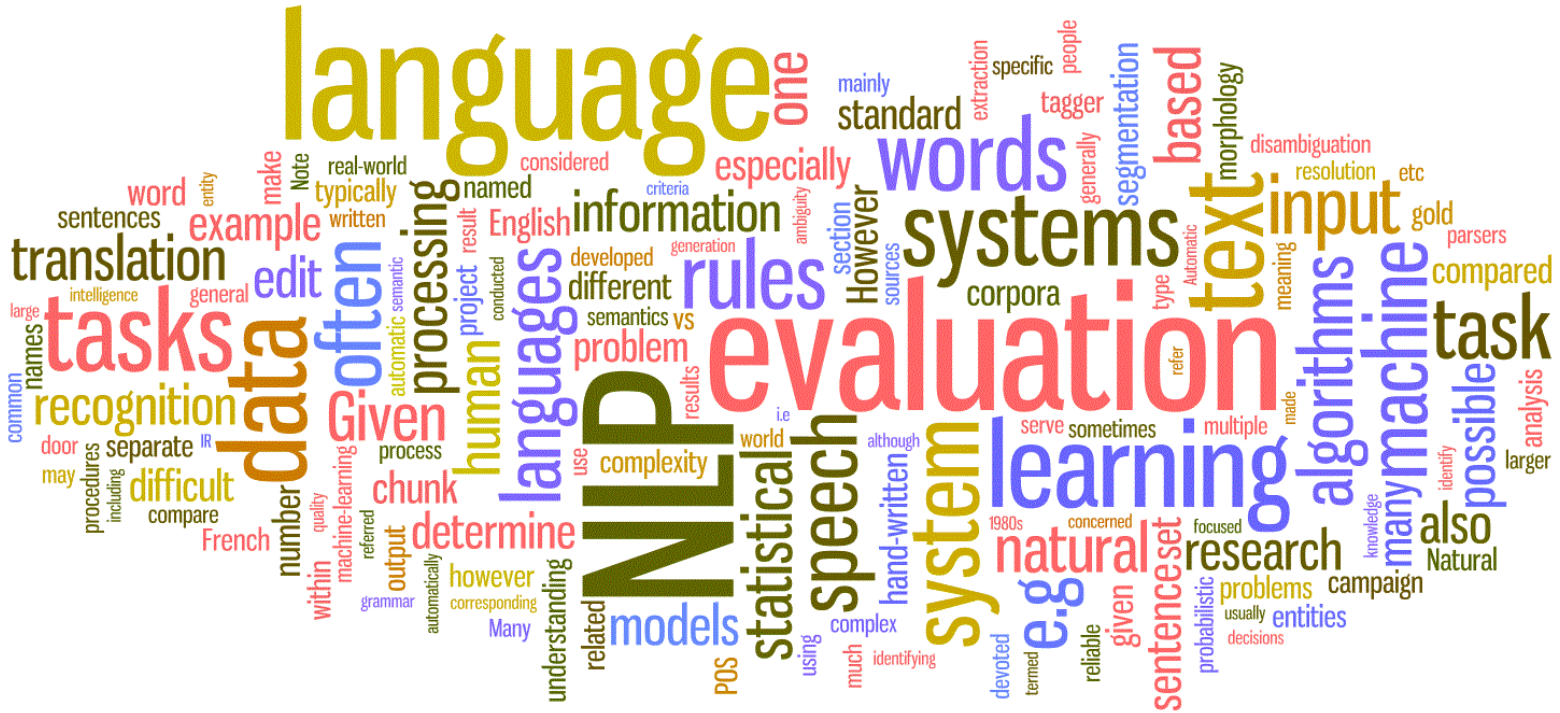

WPI

Course plan

- Original
  - Case study 3 out 3/16
  - Case study 3 due 4/6
  - Case study 4 out 4/13
  - Case study 4 due 4/27
  - Final exam 4/27

- Possible
  - Case study 3 out 3/16
  - Case study 3 due 3/30
  - Case study 4 out 4/6
  - Case study 4 due 4/20
  - Final exam 4/27

Learning **objectives** for this machine learning class.

- Supervised Regression
  - Linear Regression
  - High dimensional and non-linear
  - Model selection
  - Ridge Regression
  - **Lasso Regression**

- Advanced techniques and unsupervised learning.
  - Trees
  - Ensemble learning
  - K-means
  - **Manifold learning**

- Learn some Python packages, including:
  - scikit-learn
  - mayavi

Inspired by: http://scikit-learn.org/stable/tutorial/statishttp://scikit-learn.org/stable/tutorial/statistical_inference/supervised_learning.htmltical_inference/supervised_learning.html

WPI

Review!

The kinds of machine learning

Supervised Classification

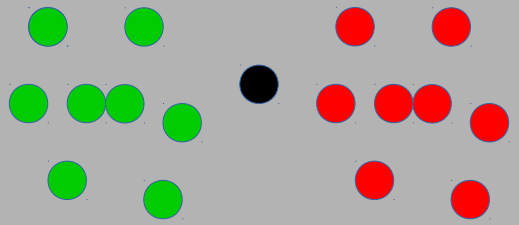

Supervised Regression

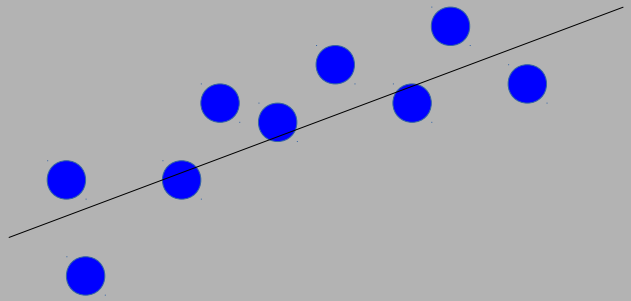

Unsupervised Clustering

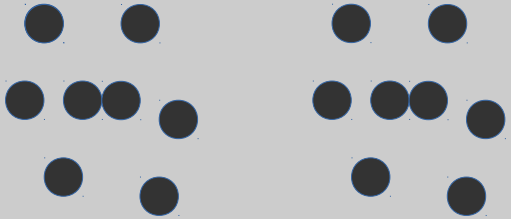

Dimension Reduction

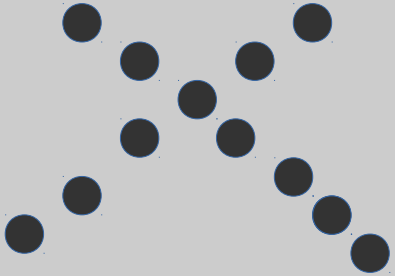

WPI

Iris data set

Features:

sepal length (cm)
sepal width (cm)
petal length (cm)
petal width (cm)

"Iris virginica" by Frank Mayfield - originally posted to Flickr as Iris virginica shrevei BLUE FLAG. Licensed under Creative Commons Attribution-Share Alike 2.0 via Wikimedia Commons - http://commons.wikimedia.org/wiki/File:Iris_virginica.jpg#mediaviewer/File:Iris_virginica.jpg

Catagories:

setosa
versicolor
virginica

"Kosaciec szczecinkowaty Iris setosa". Licensed under Creative Commons Attribution-Share Alike 3.0 via Wikimedia Commons - http://commons.wikimedia.org/wiki/File:Kosaciec_szczecinkowaty_Iris_setosa.jpg#mediaviewer/File:Kosaciec_szczecinkowaty_Iris_setosa.jpg

"Iris versicolor 3". Licensed under Creative Commons Attribution-Share Alike 3.0 via Wikimedia Commons - http://commons.wikimedia.org/wiki/File:Iris_versicolor_3.jpg#mediaviewer/File:Iris_versicolor_3.jpg

WPI

Iris data set

$f_1(x, y, z, w)$

$f_2(x, y, z, w)$

$f_3(x, y, z, w)$

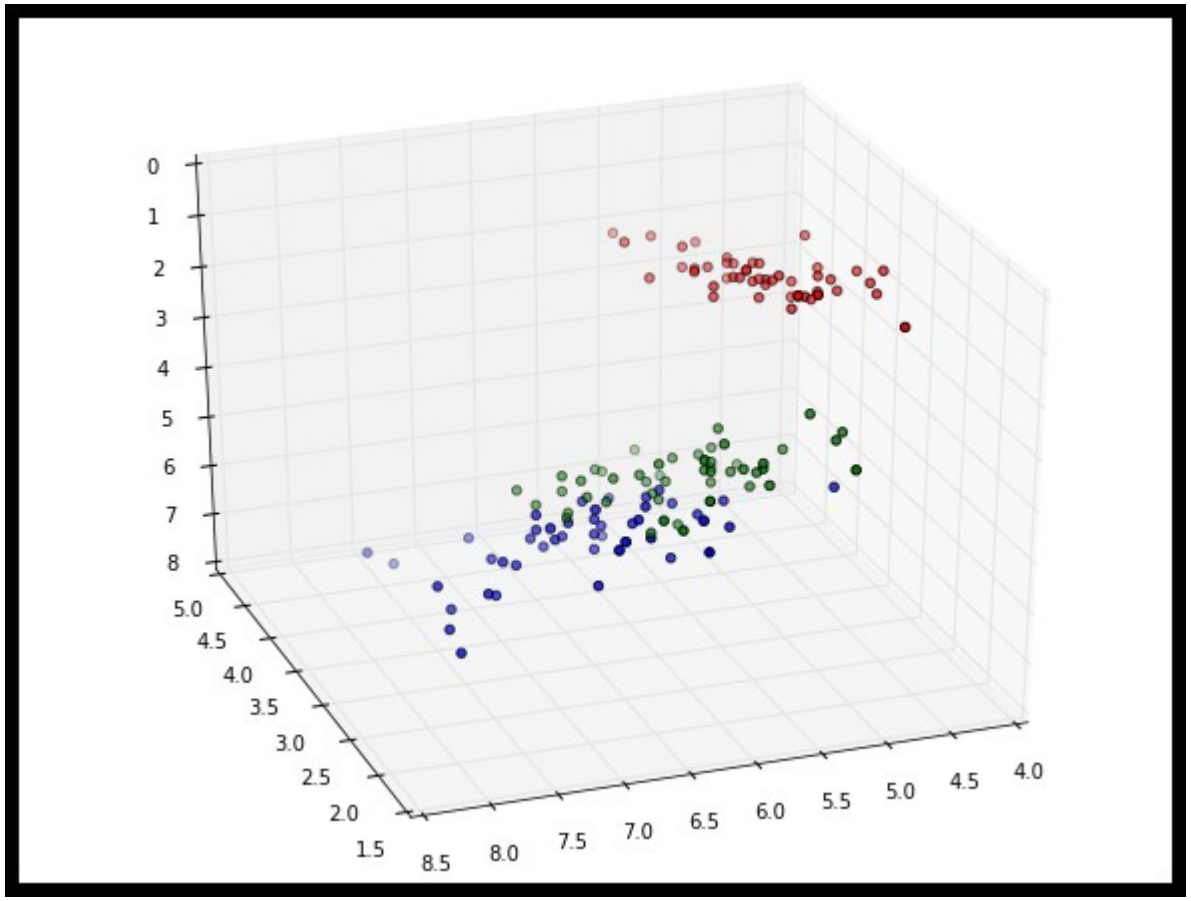

WPI

What is PCA?

- Principle Component Analysis
  - Commonly used tool for visualization and data pre-processing.

Linear fn

$z_1 = a_1 x_1 + b_1 y_1 + c_1 z_1 + d_1 w_1$

Pick a_1, b_1, c_1, d_1 to maximize variance

WPI

What is Linear Support Vector Machine (SVM)?

- Maximum margin classifier
  - Computes a linear “decision boundary” that splits the data into two regions.
  - Allows one to predict a classification of a point based upon which side of the decision boundary it lay on.

SVM

as far as possible from the closest points

maximum margin

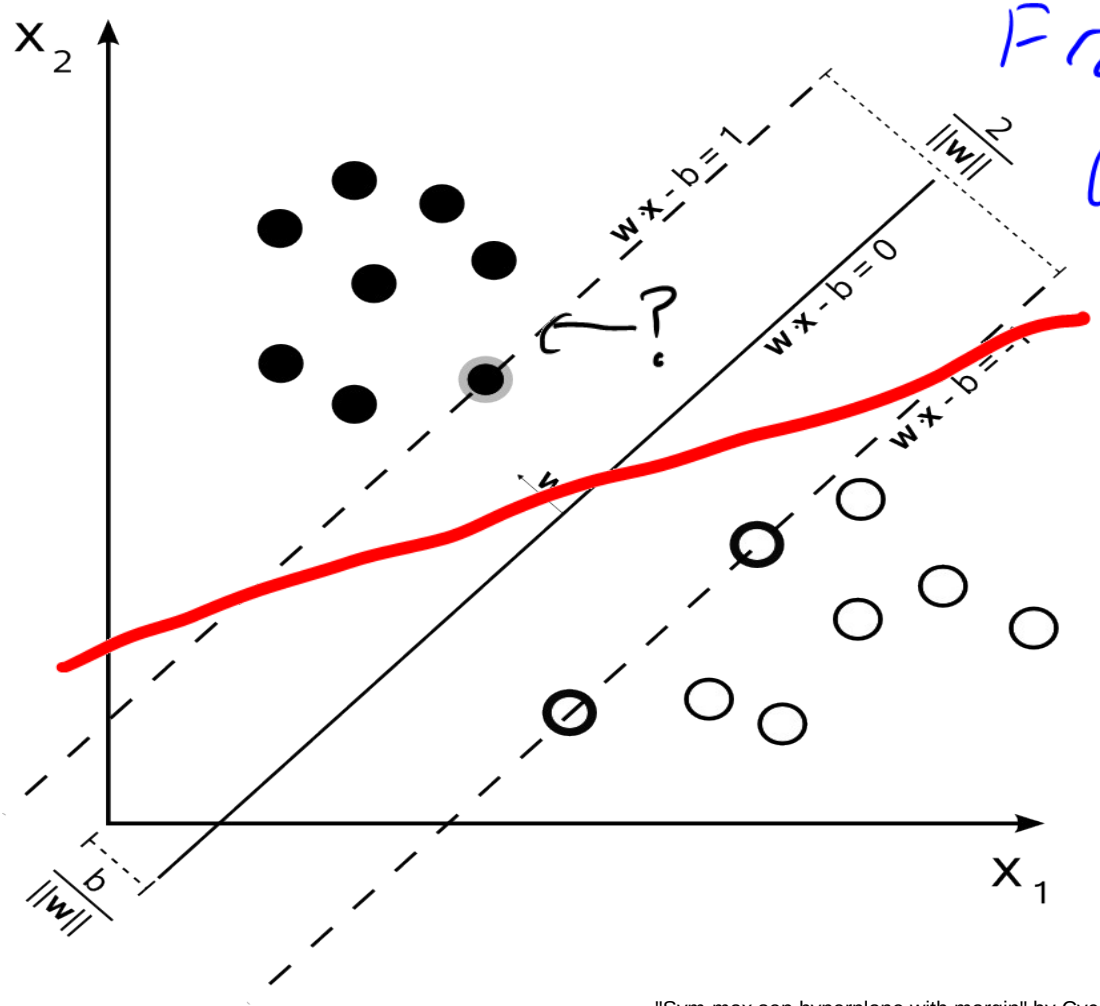

WPI

SVM

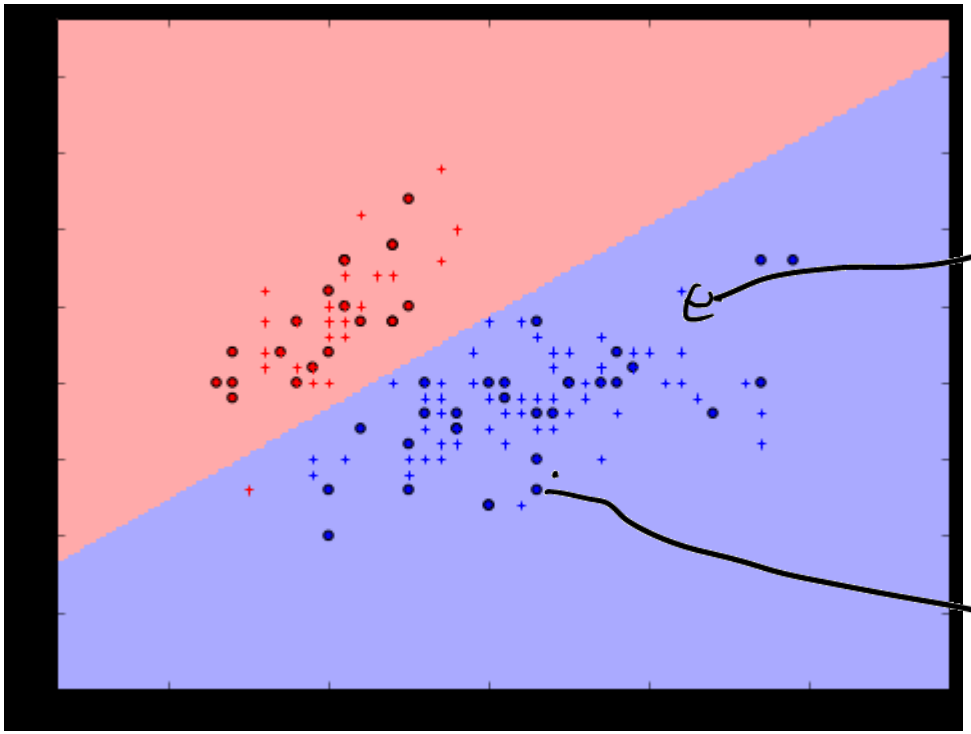

But wait! Training vs. testing!

Testing data

computing errors

Train data

Computing your model

What is K-NN?

- K-nearest neighbors
- Another common classification algorithm
  - Perhaps the most common

K-NN

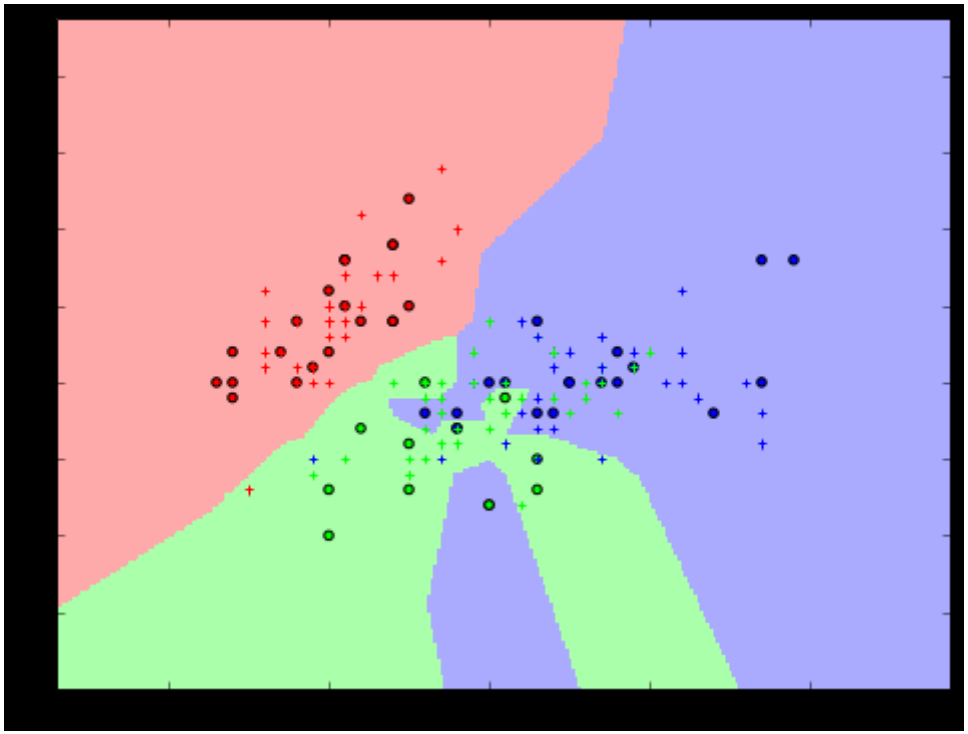

WPI

New Material!

scikit-learn

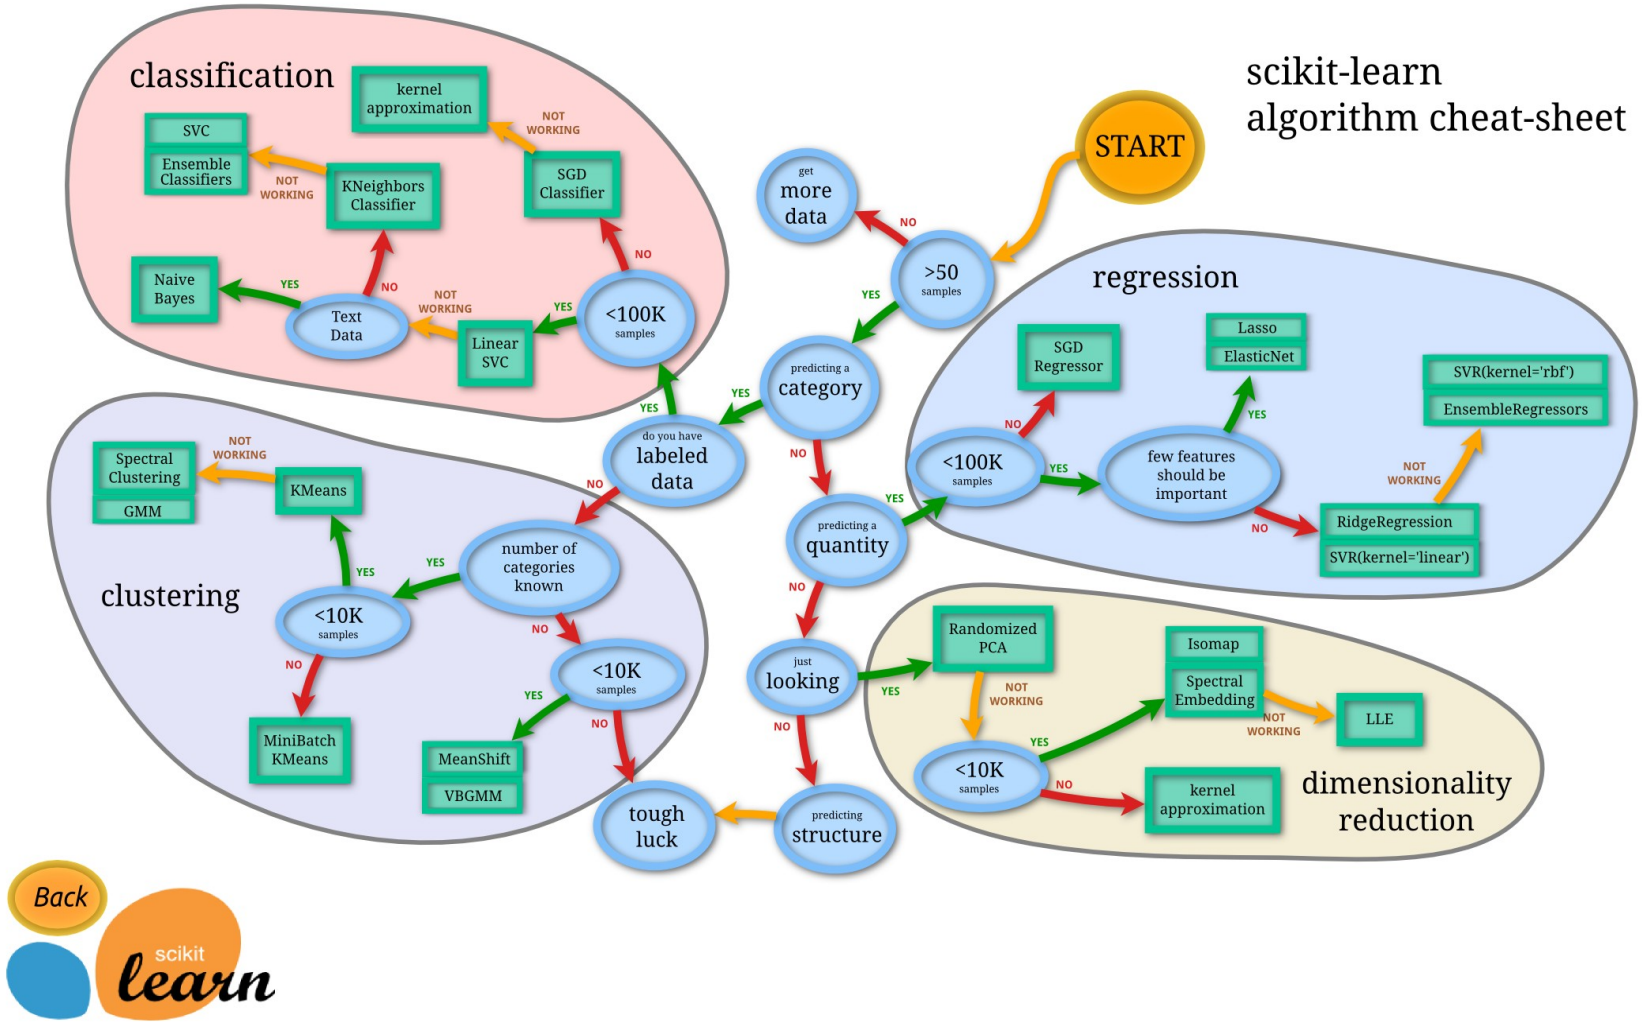

http://scikit-learn.org/stable/tutorial/machine_learning_map/index.html

WPI

Linear Regression flow chart

Fit

$$RSS = (y_1 - \hat{y}_1)^2 + \cdots + (y_p - \hat{y}_p)^2$$

Compute minimum

$$\hat{\boldsymbol{\beta}} = (X^T X)^{-1} X^T y$$

Predict

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X$$

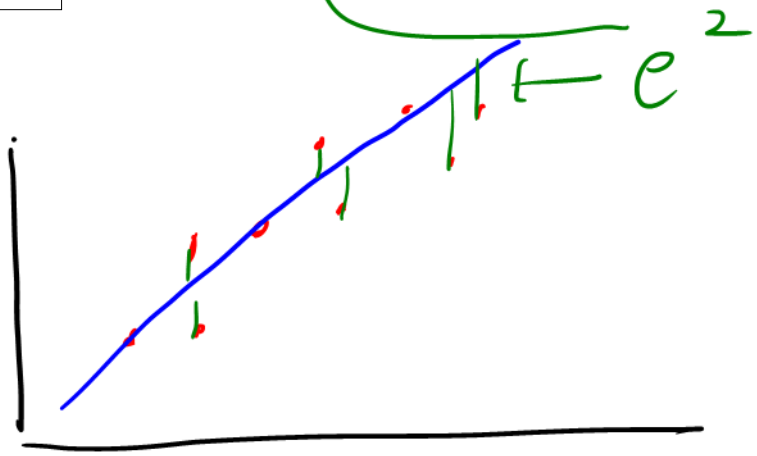

Error

$$RSE = \sqrt{\frac{RSS}{n - (p+1)}}$$

Plug in data (testing and *training*) to see how we did.

WPI

Relationship between Linear Regression and PCA...

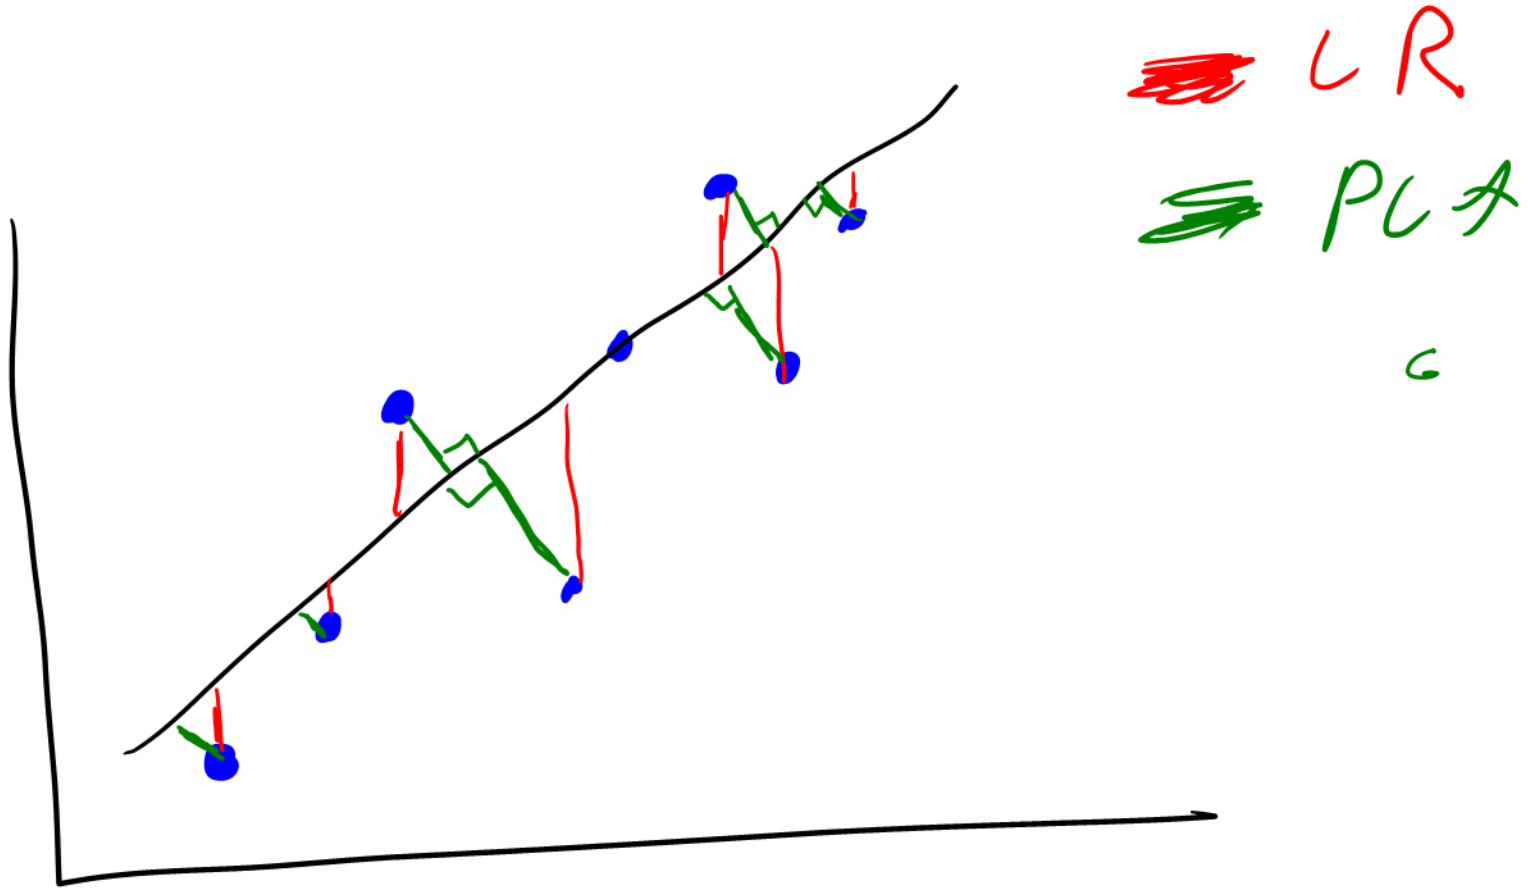

WPI

Multiple-linear regression

gas mileage ↓ weight ↓ engine size ↓

$$Y = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k$$

$$y_i = \beta_0 + \beta_1 x_{1,i} + \cdots + \beta_1 x_{k,i} \quad \hat{y}_0 = \beta_0 + \beta_1 x_{1,0} + \cdots + \beta_1 x_{k,0}$$

“Non-linear” regression

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1 X_2$$

See it in Python

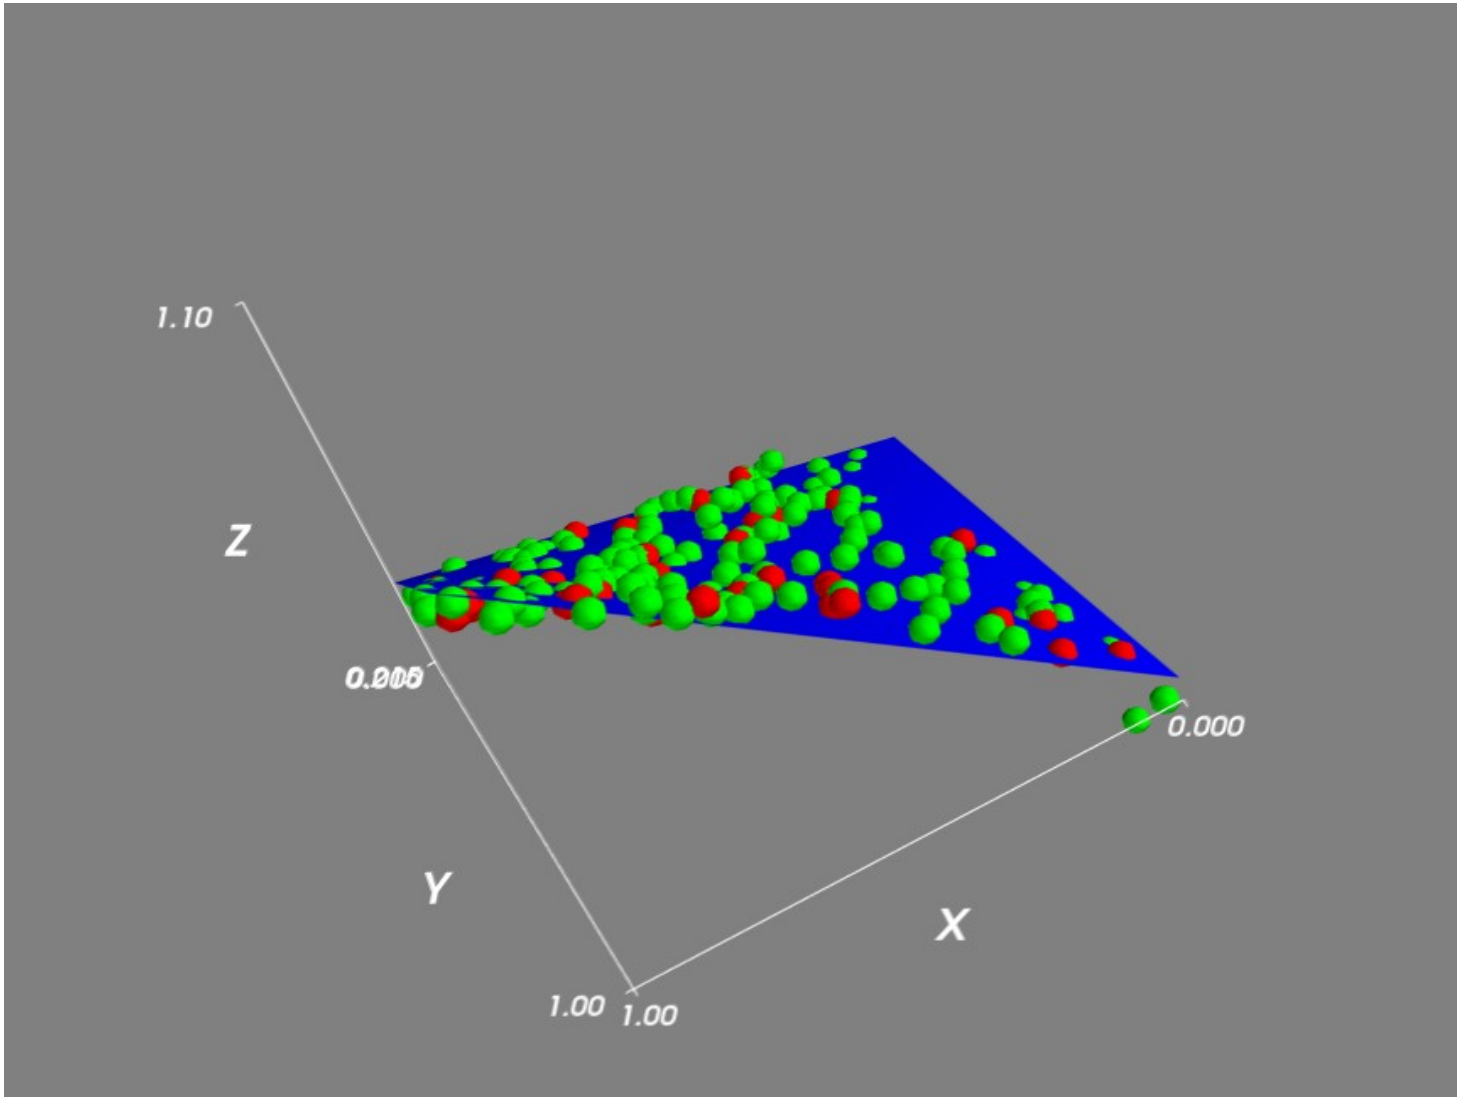

Too much of a good thing...

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_1 + \beta_3 X_1 X_2$$

$+\beta_4 X_1^2 X_2^2 \quad +\beta_5 X_1^3 X_2^7 \quad +\beta_6 \cos(X_1) e^{X_2}$

WPI

See it in Python

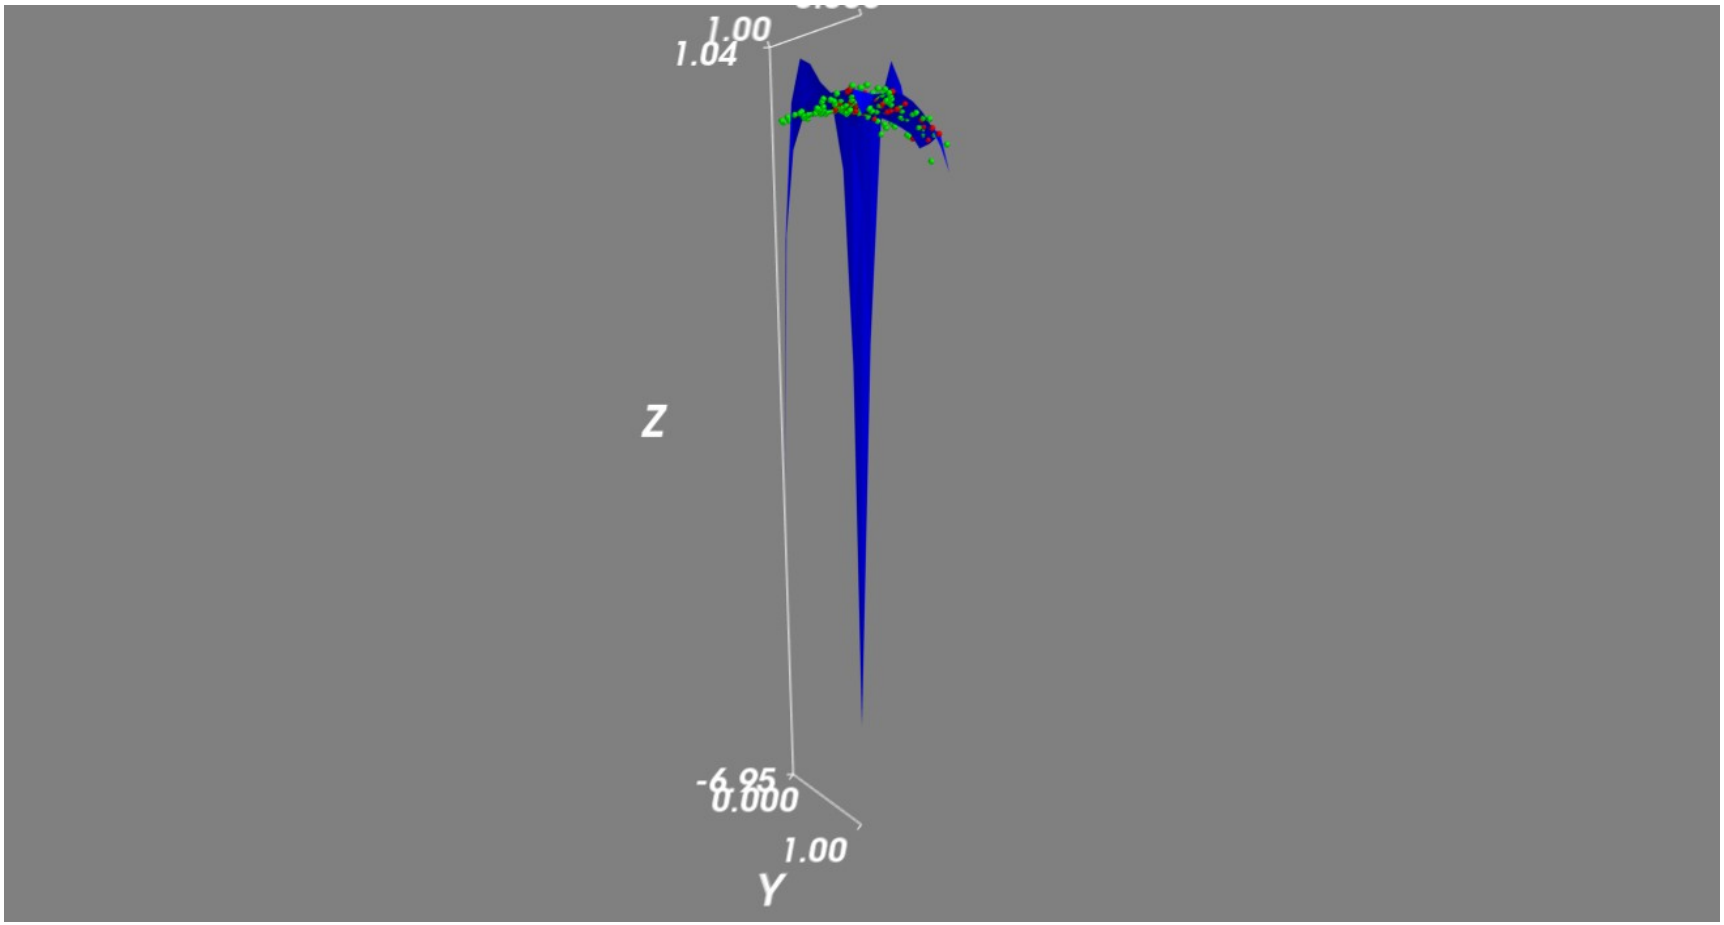

WPI

Cross Validation

- Validation set

- K-fold
- Leave-one-out cross validation (LOOCV)

4-Fold Cross validation

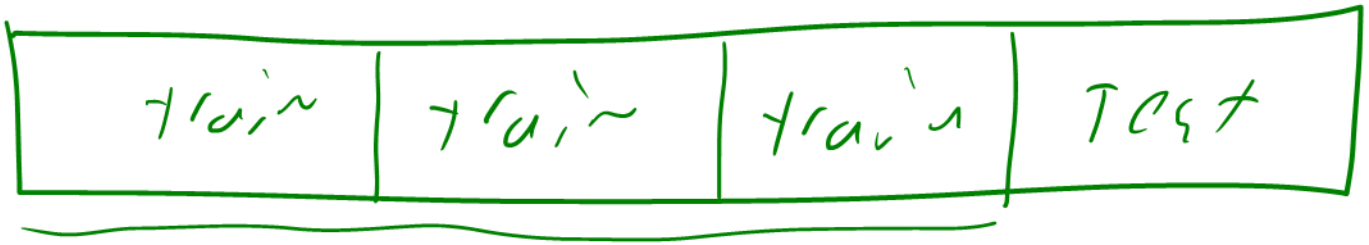

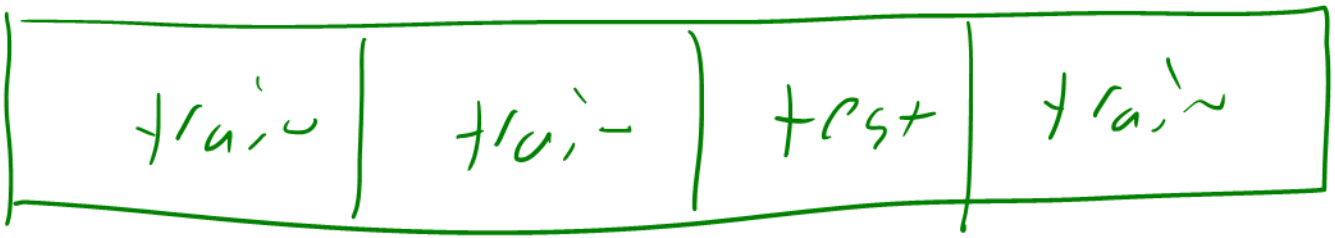

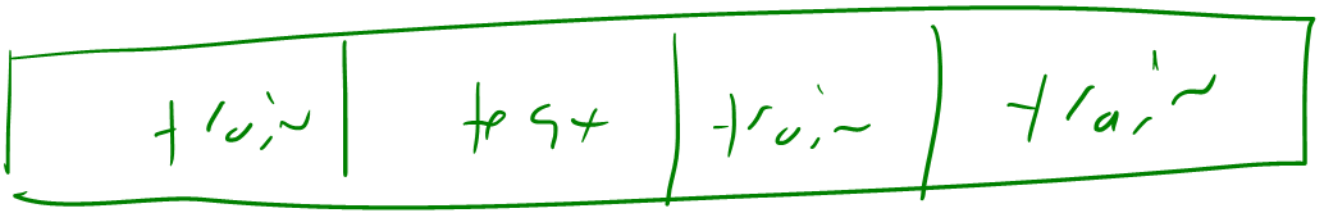

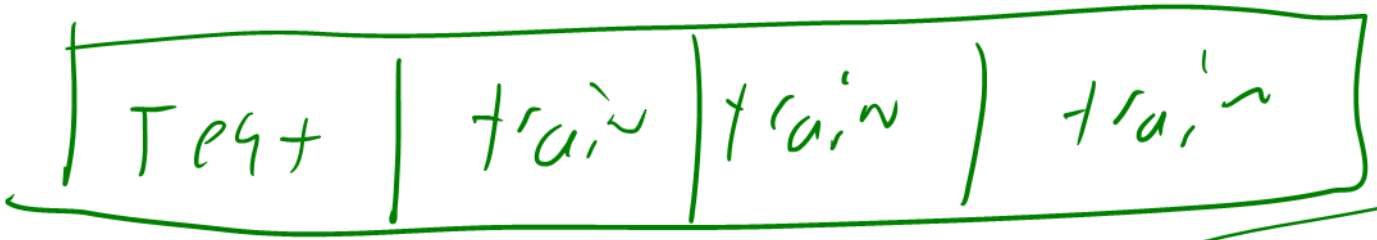

L O O C V

Leave One Out Cross Validation

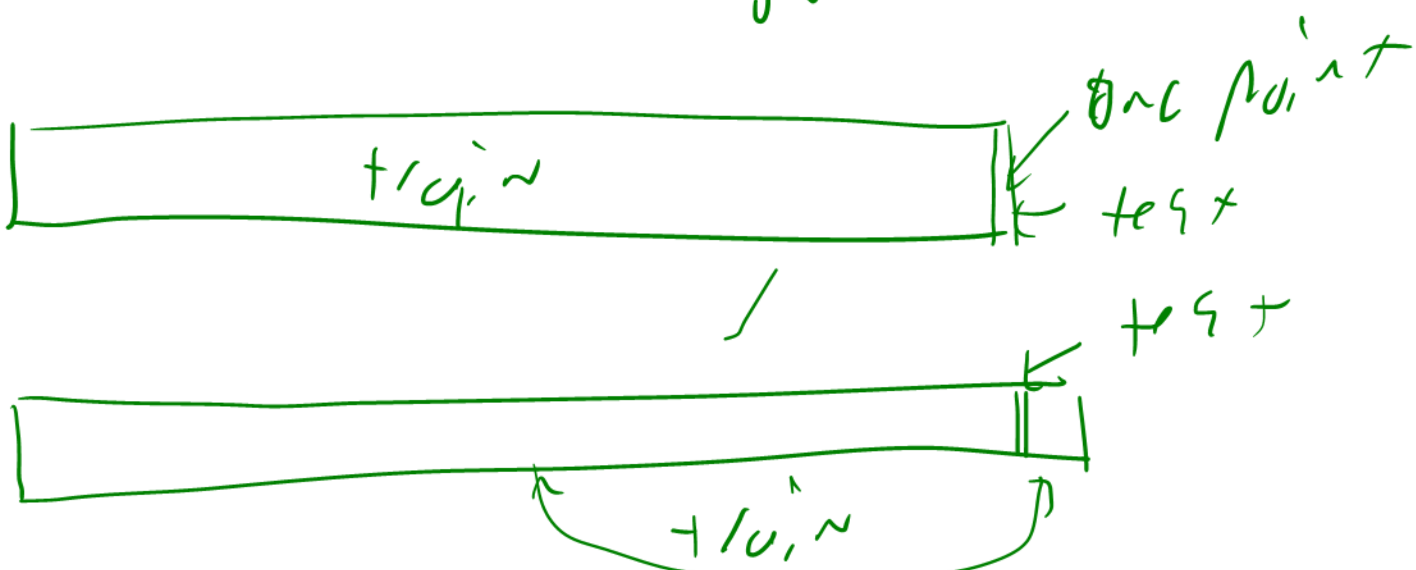

Feature selection

- Can someone describe:
  - Best-subset selection
  - Forward stepwise selection
  - Backward stepwise selection
- Recursive Feature Elimination (RFE) is what we will use today: http://axon.cs.byu.edu/Dan/778/papers/Feature%20Selection/guyon*.pdf
  - It would take us too far astray to talk about the details of this algorithm, but it is a close cousin of backward selection.
  - Steps
    - 1. Train the classifier.
    - 2. Compute the ranking criterion for all features.
    - 3. Remove the feature with smallest ranking criterion.

See it in Python

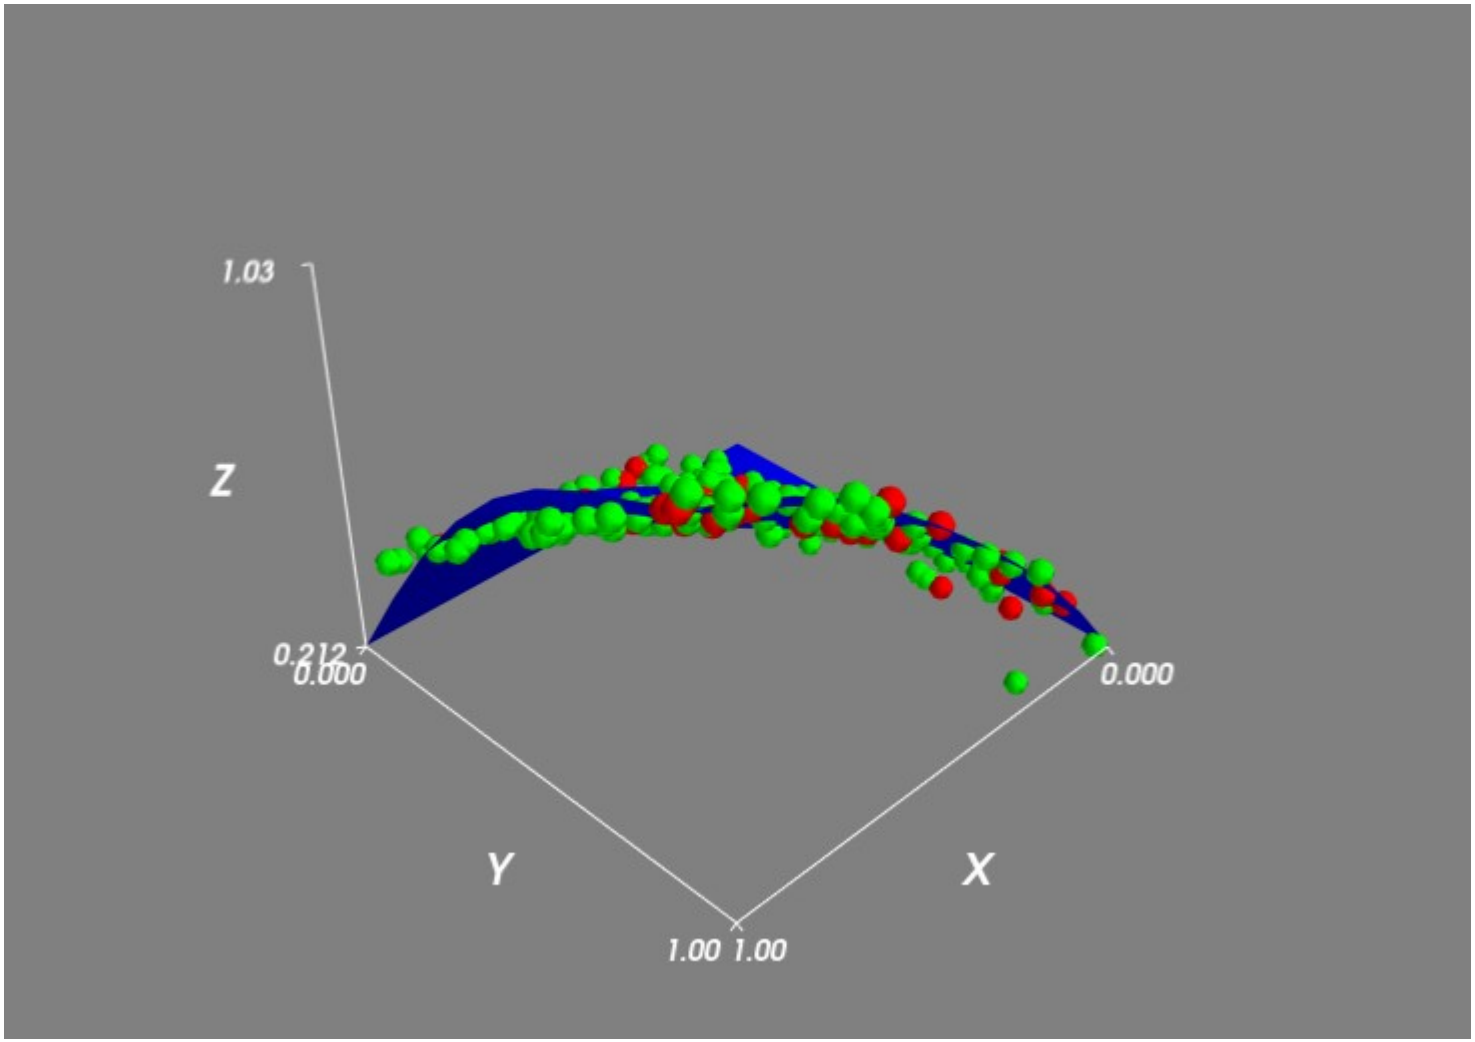

WPI

Ridge Regression

$$\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_1)^2 + \lambda(\beta_0^2 + \beta_1^2) = RSS + \lambda(\beta_0^2 + \beta_1^2)$$

$$\min_{\beta}\left[\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_1)^2\right] \text{ s.t.} \beta_0^2 + \beta_1^2 \leq s$$

Budget

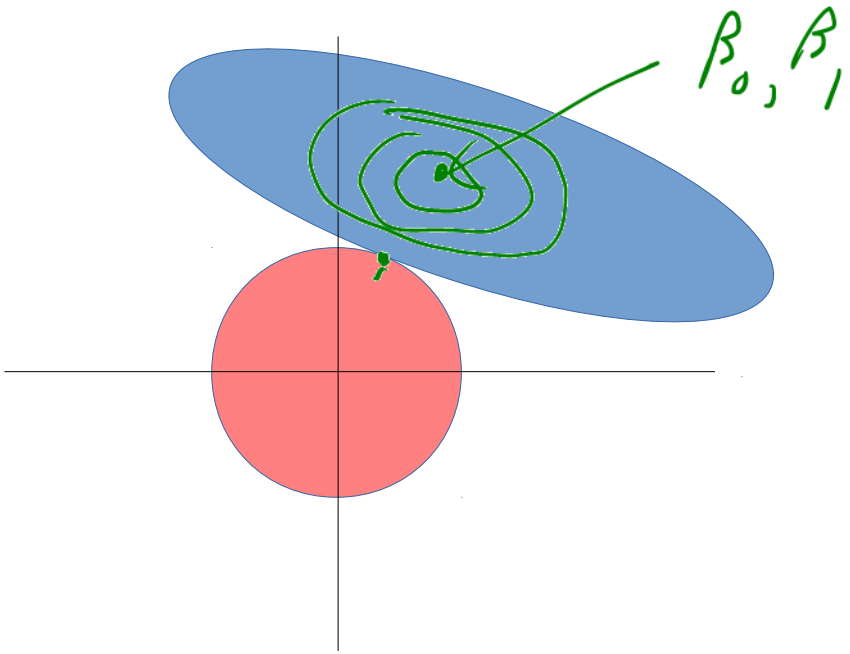

WPI

See it in Python

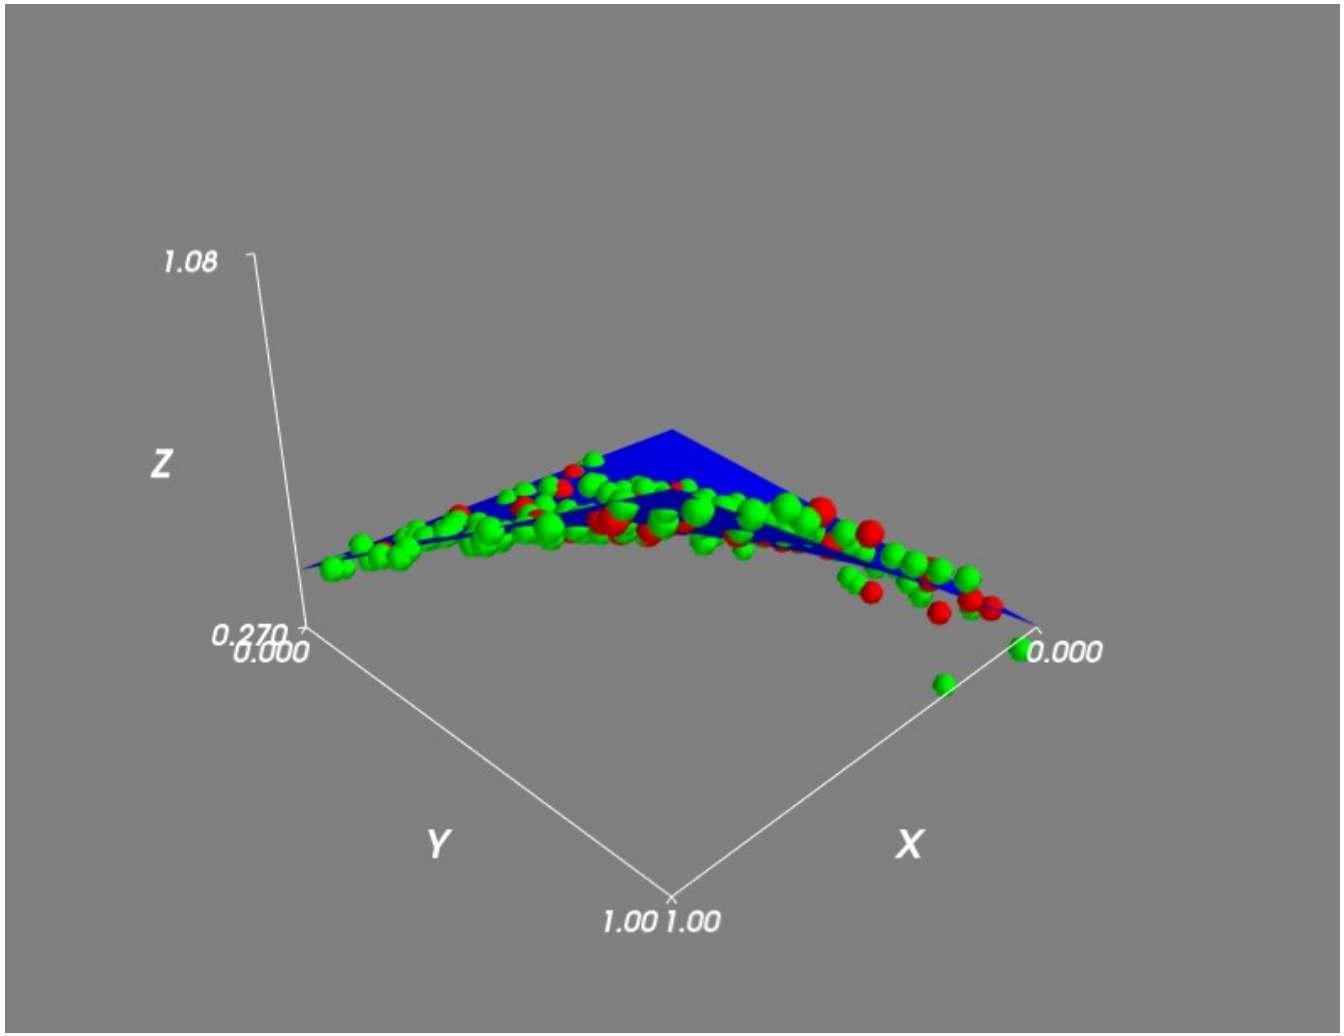

WPI

Lasso Regression

$$\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_1)^2 + \lambda(|\beta_0| + |\beta_1|) = RSS + \lambda(|\beta_0| + |\beta_1|)$$

$$\min_{\beta}\left[\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_1)^2\right] \text{s.t.} |\beta_0| + |\beta_1| \leq s$$

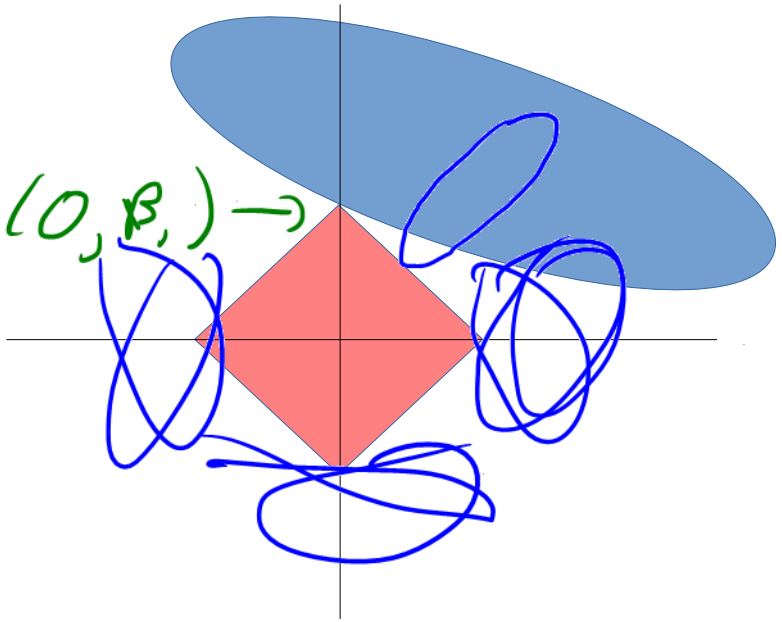

WPI

See it in Python

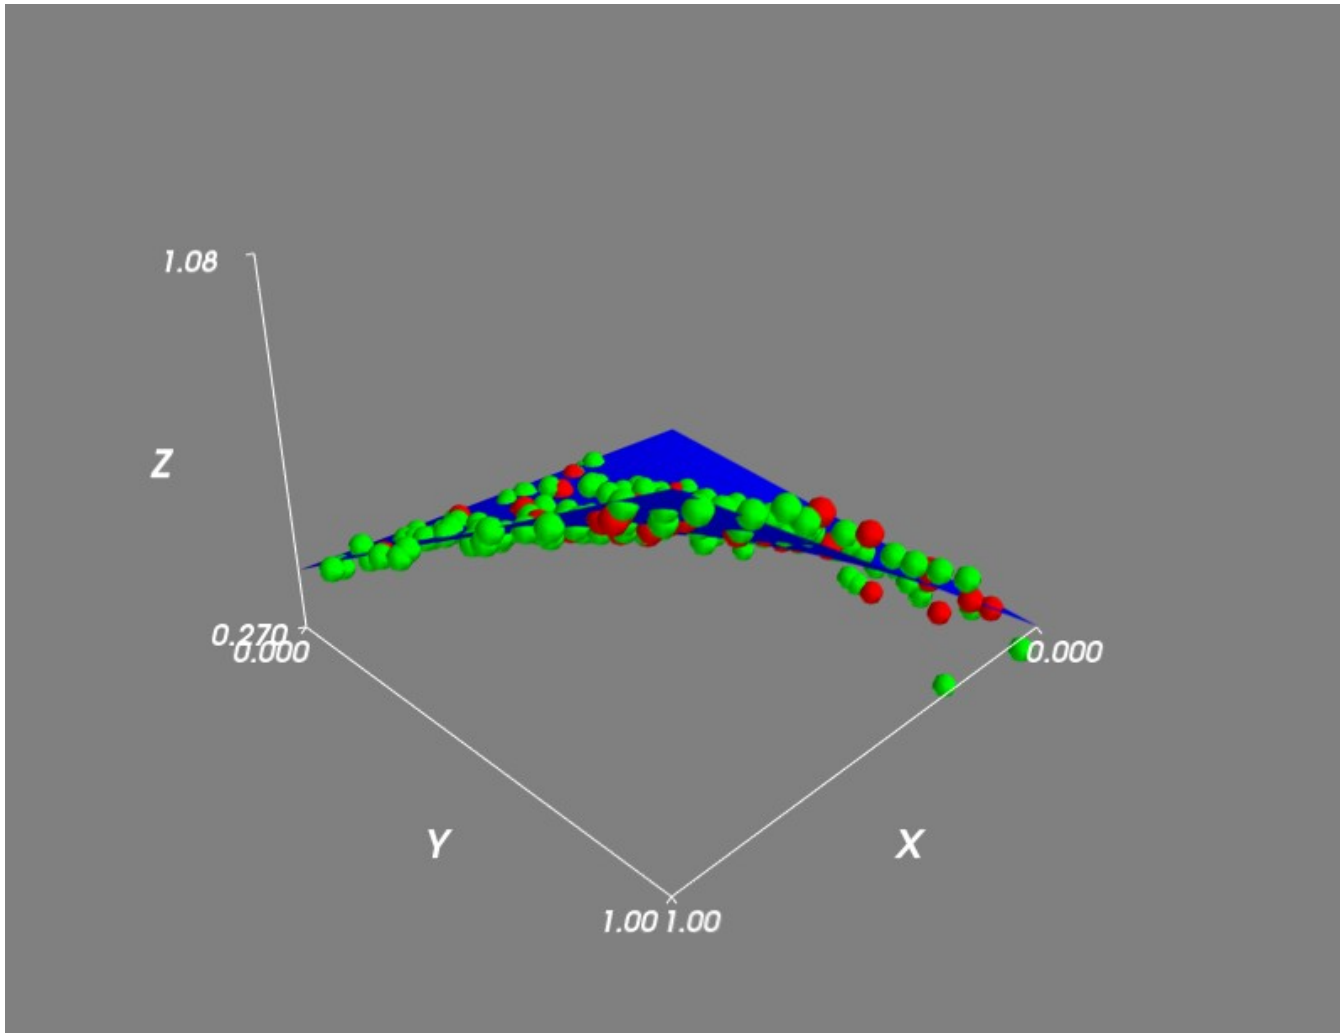

WPI

Let's move on...

Supervised Classification

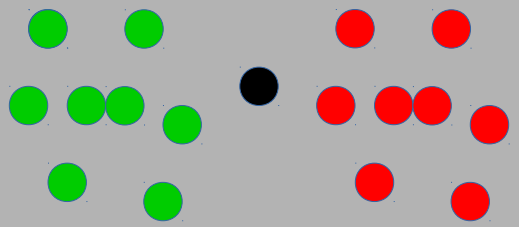

Supervised Regression

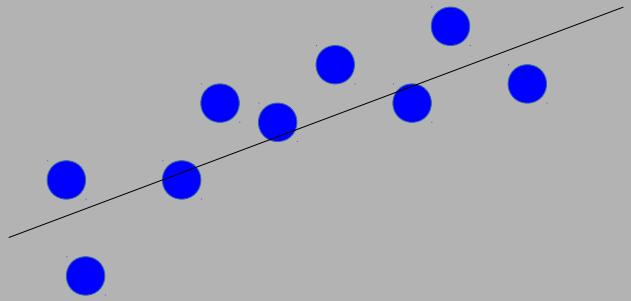

Unsupervised Clustering

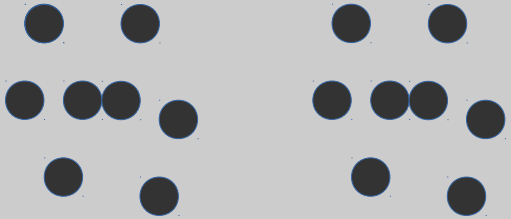

Dimension Reduction

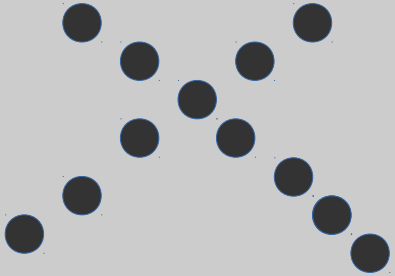

WPI

scikit-learn

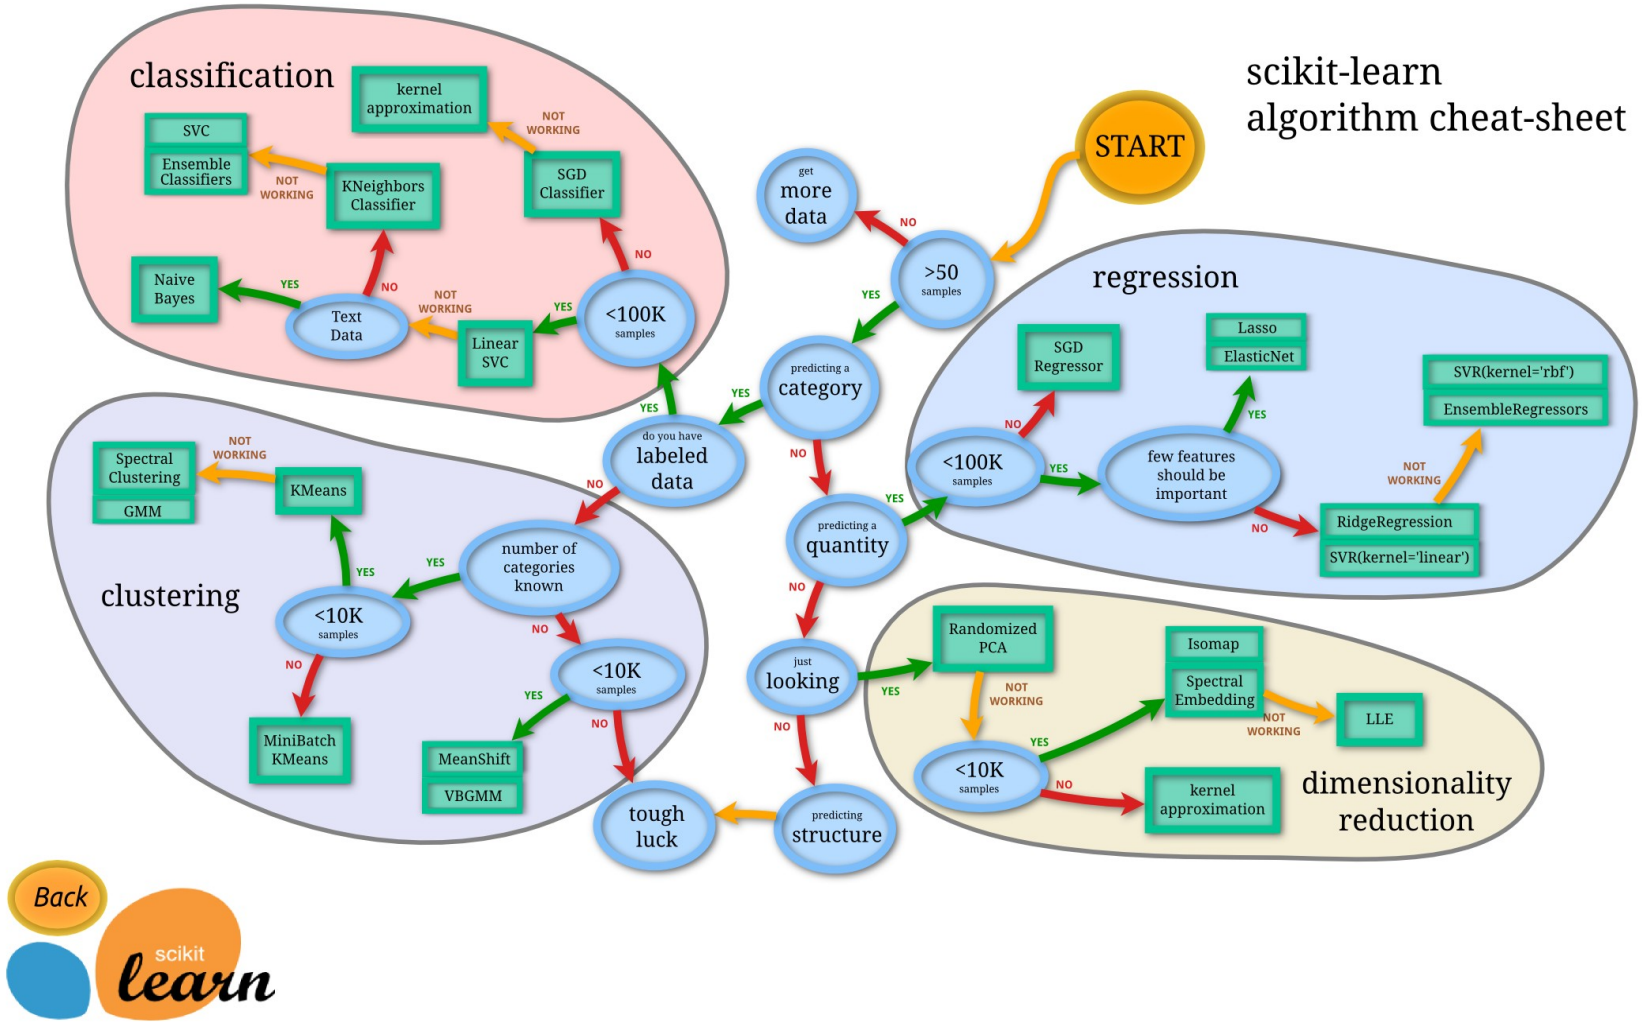

http://scikit-learn.org/stable/tutorial/machine_learning_map/index.html

WPI

Decision tree

- Quite commonly used in data mining.
- Each node in the tree splits the data into (classically) two groups.
  - To make things easy (and fast) you classically perform each split on a single variable.
- Each leaf node then represents a value (or perhaps range of values) for the response based upon the input variables.

Titanic data

Probability of survival
Percentage in that leaf

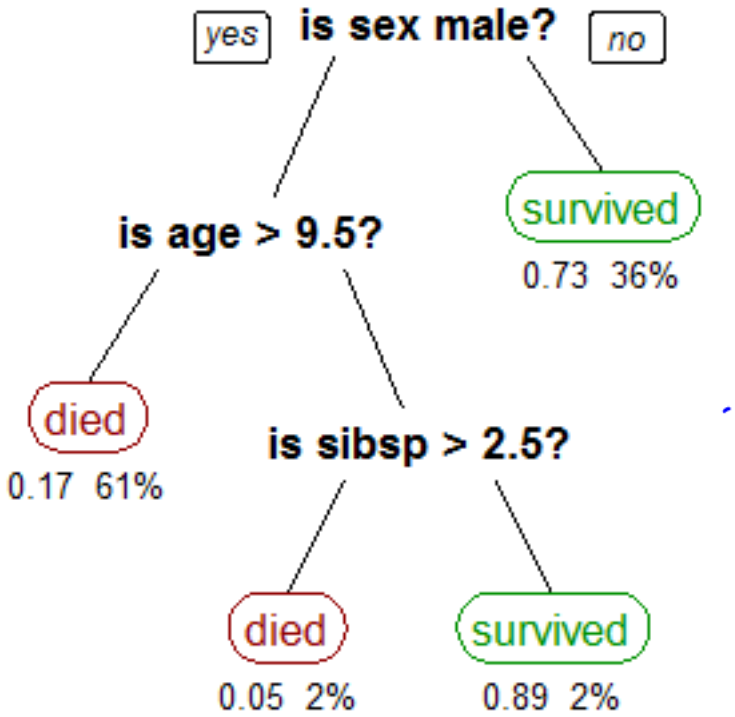

WPI

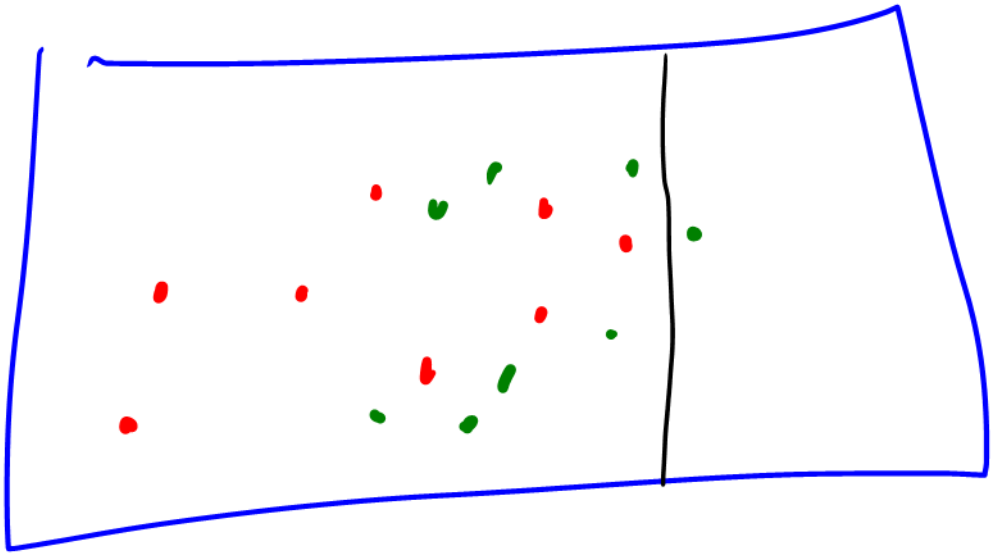

Decision tree: Making the splits

$$\hat{p}_{mk} = Pr(Y = k | X \text{ is in region } k)$$

$$E = \max_k \hat{p}_{mk}$$

Entropy

$$G = \sum_{k=1}^{K} \hat{p}_{mk}(1 - \hat{p}_{mk})$$

Gini

WPI

Decision Trees

1) which subset
2) which predictor
3) where in the predictor

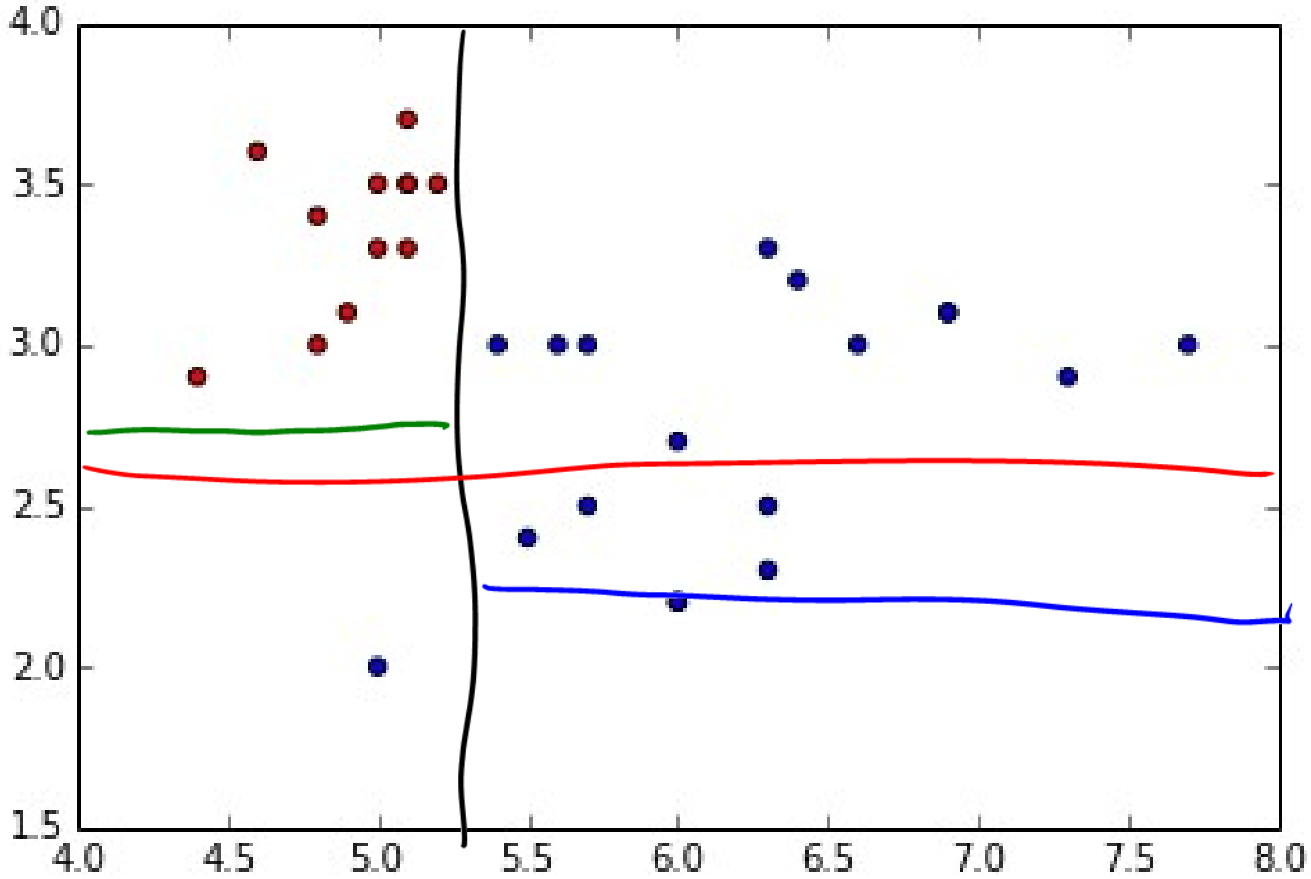

very valid
valid
invalid

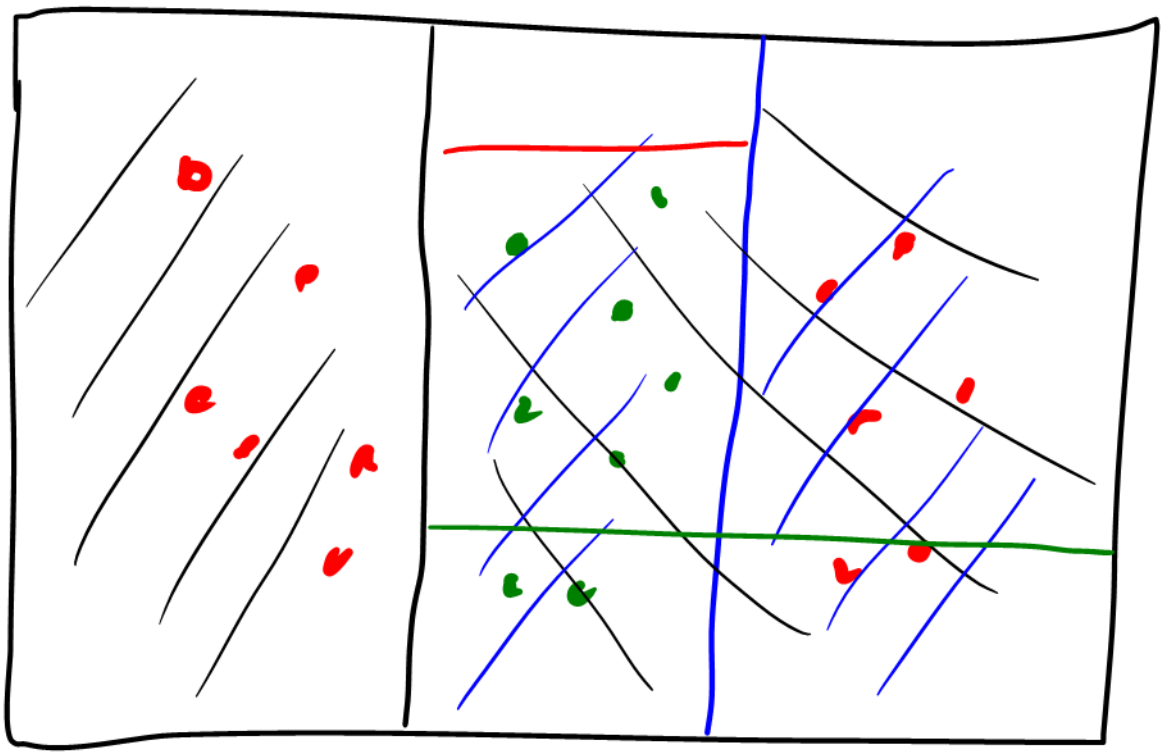

OK

OK

See it in Python

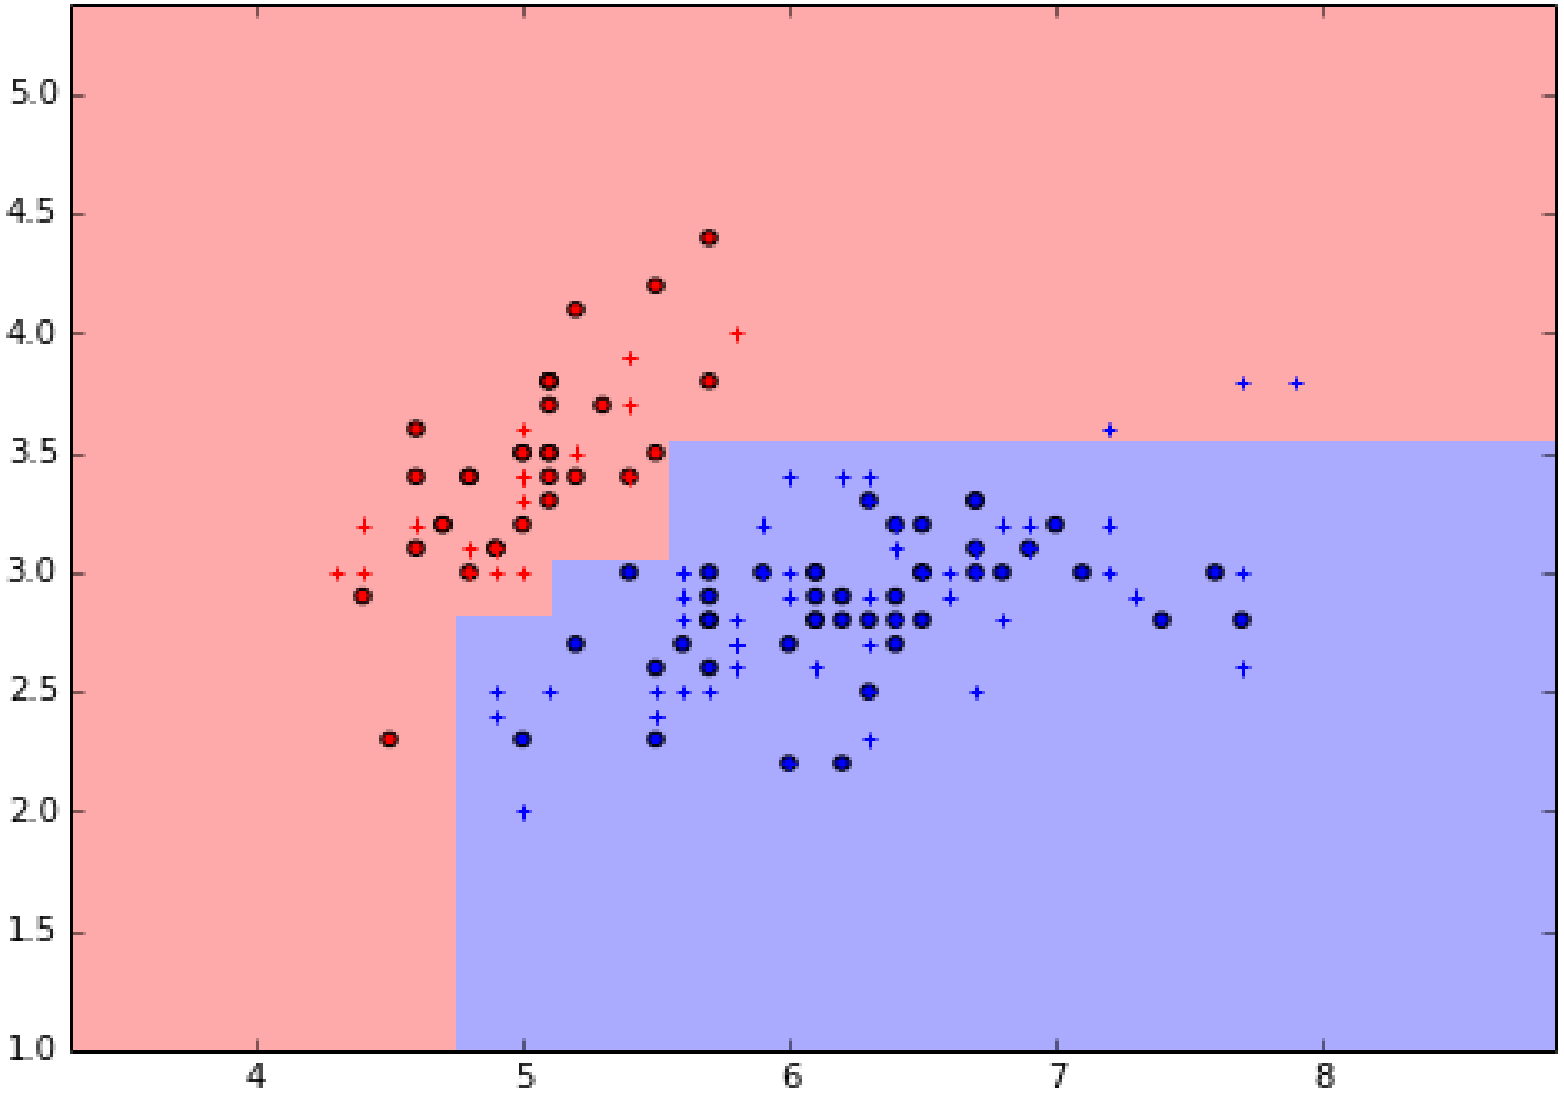

WPI

Regression trees

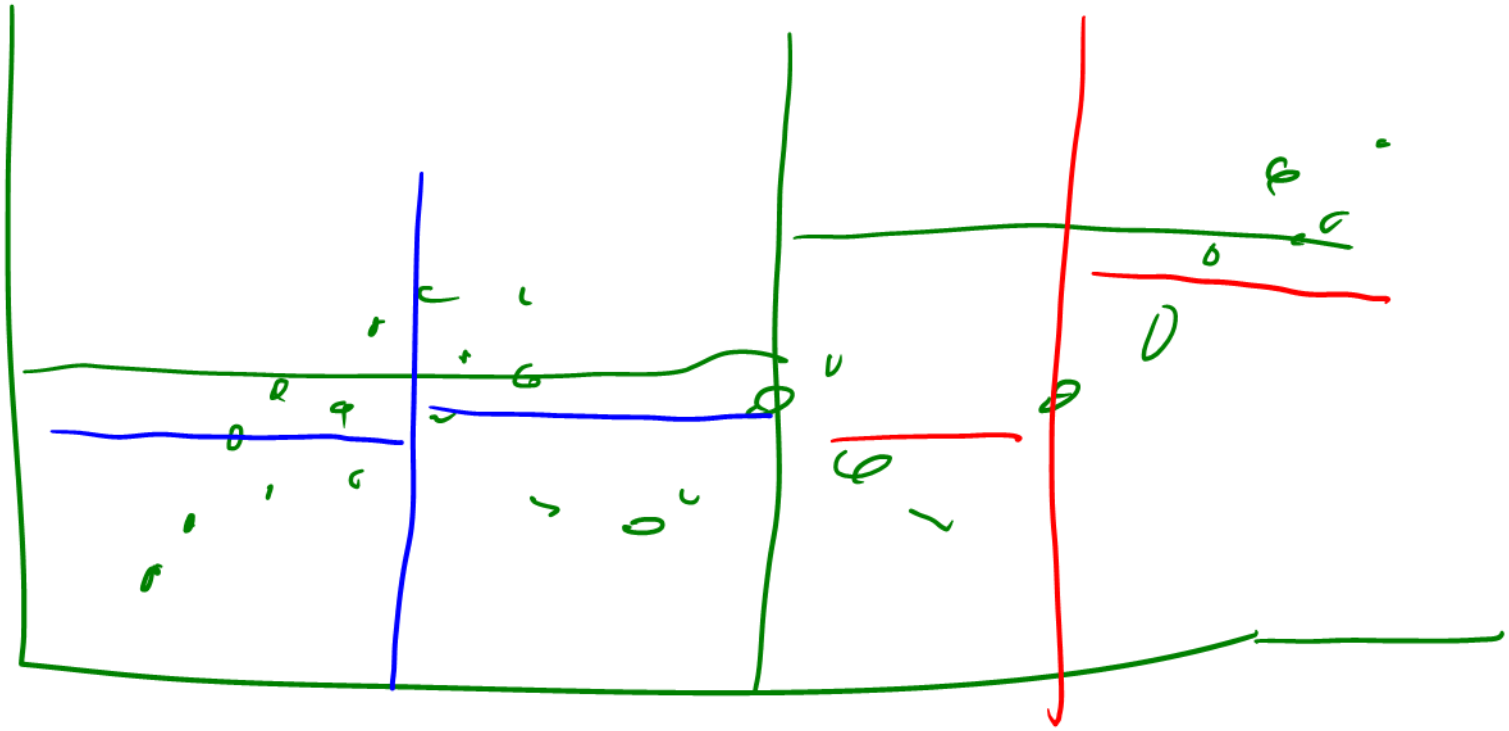

See it in Python

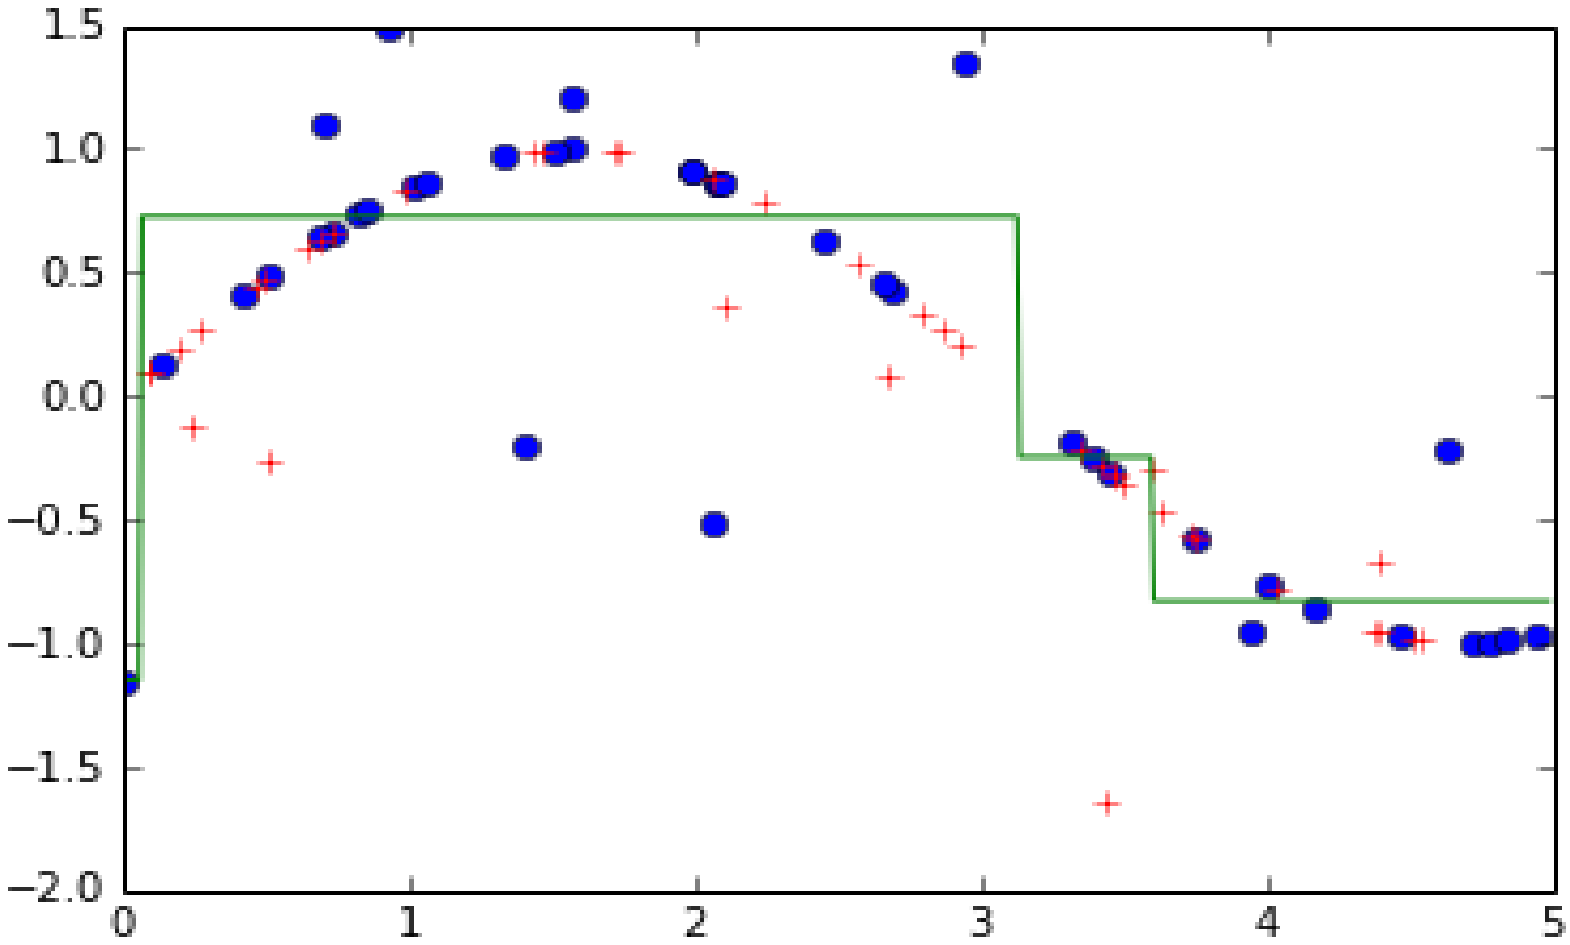

WPI

Ensemble Learning: Random Forest

- As you can tell by the name, this idea revolves around having many trees.

Tree bagging: Bootstrap aggregation

- Bootstrapping is one of my favorite algorithms in statistical learning.
- An extremely powerful idea for doing statistical learning with limited data.
- Generate many random samples of your data, with replacement, and train a tree on each...

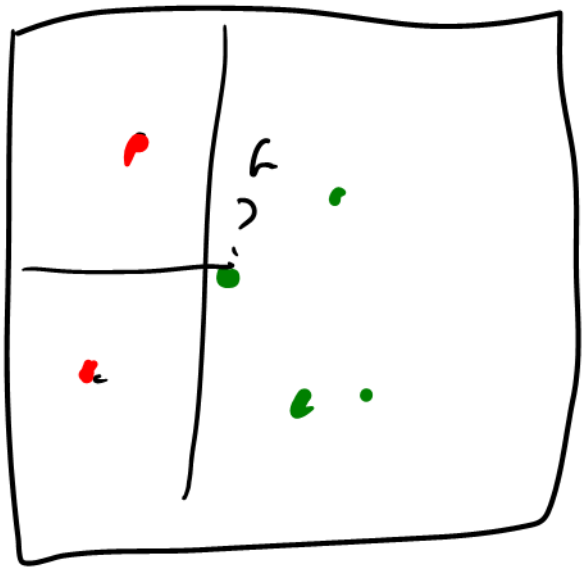

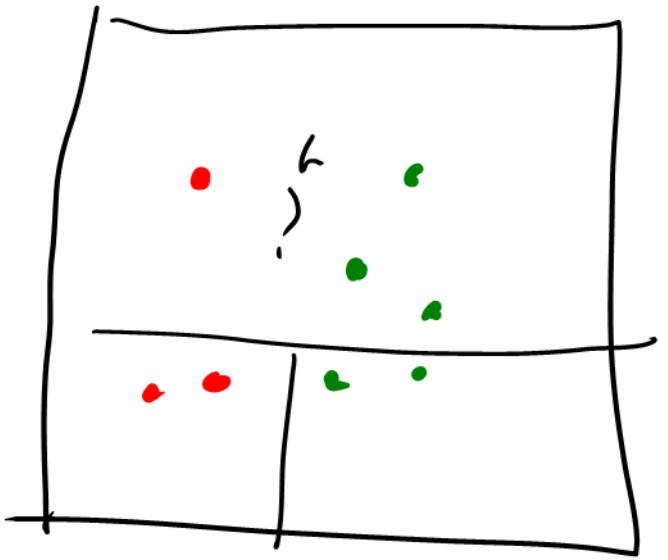

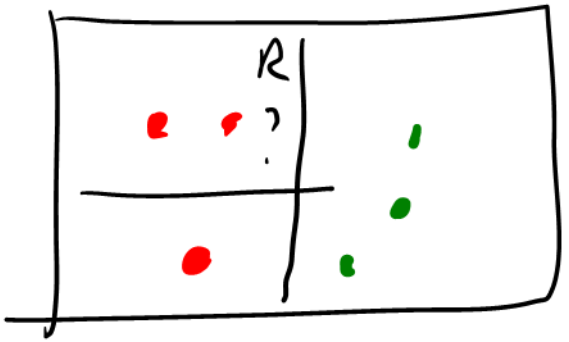

G, G, R,

answer

G

Random forests

- Add even more randomness by randomly selecting for each tree a subset of the features it is allowed to split on.
  - Reduces correlation between the trees!
  - Not all trees can pick the “obvious” best predictor to split on first.

See it in Python

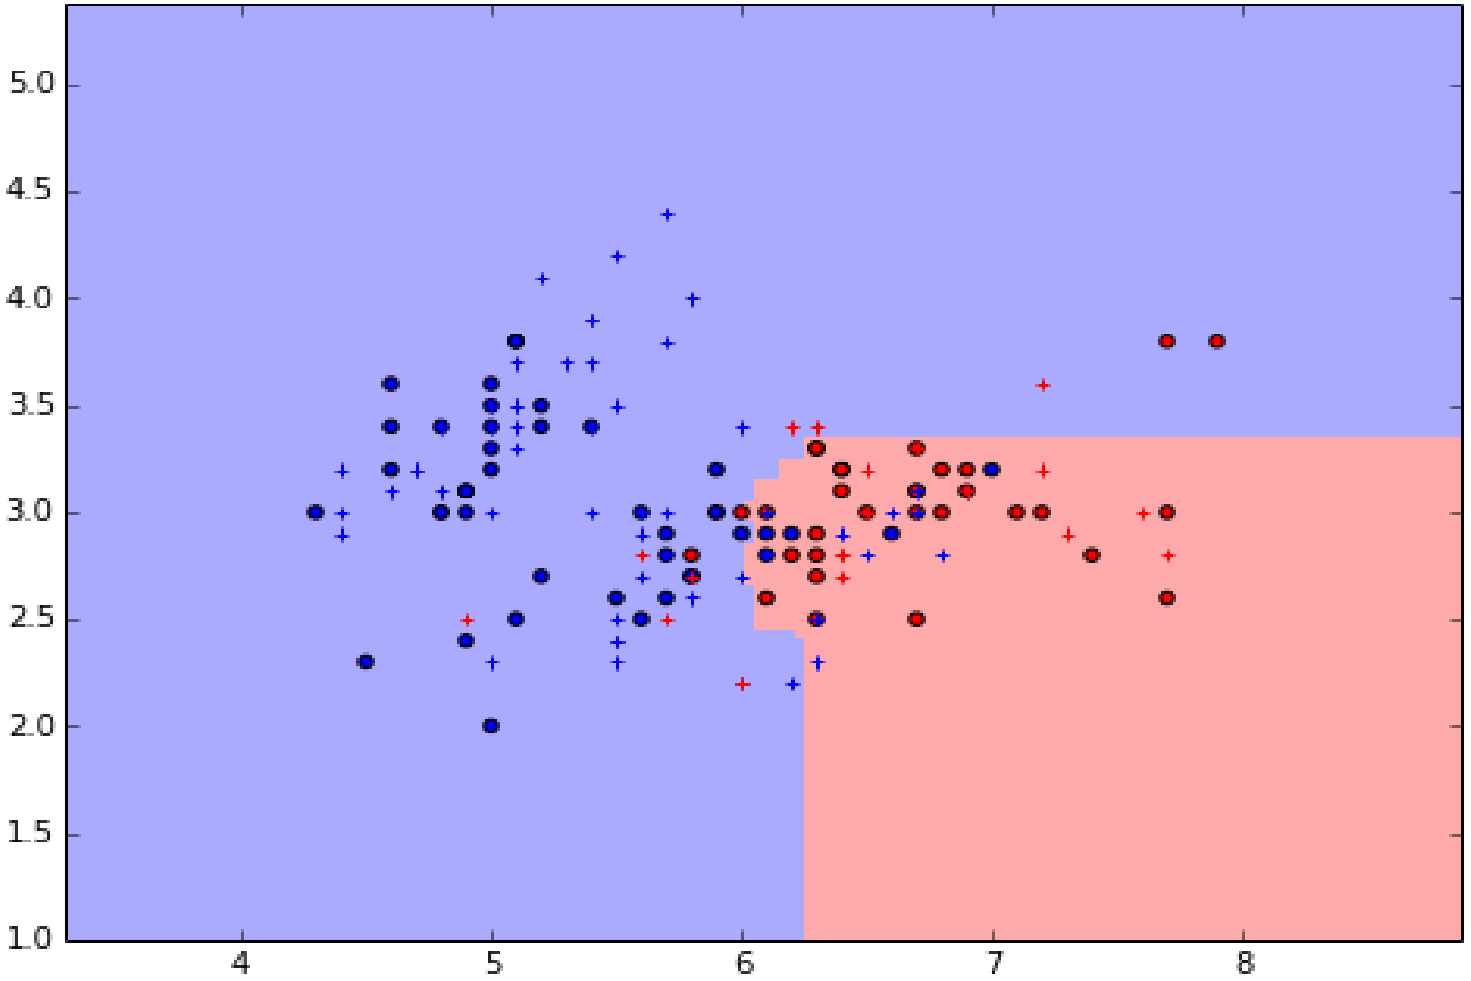

WPI

K-Means

- Perhaps the single most used unsupervised classification algorithm.
- Given a number of classes k, divide the data into groups so that the distance within a group is “small” compared to the distances between groups.

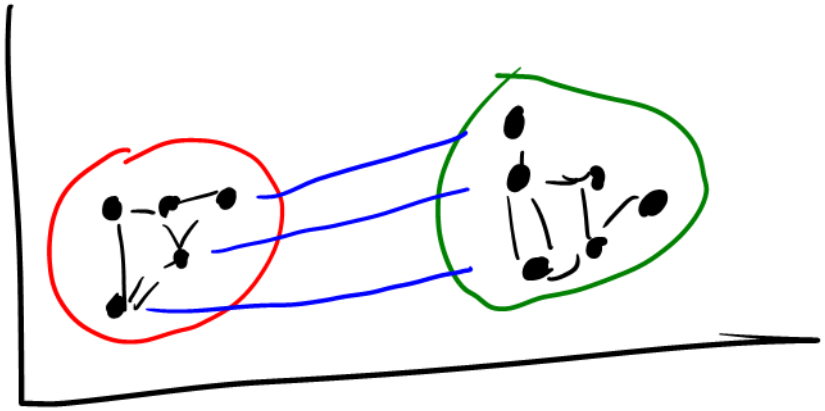

WPI

K-Means

$$\arg\min_{S} \sum_{i=1}^{k} \sum_{x \in S_i} \|x - \mu_i\|^2$$

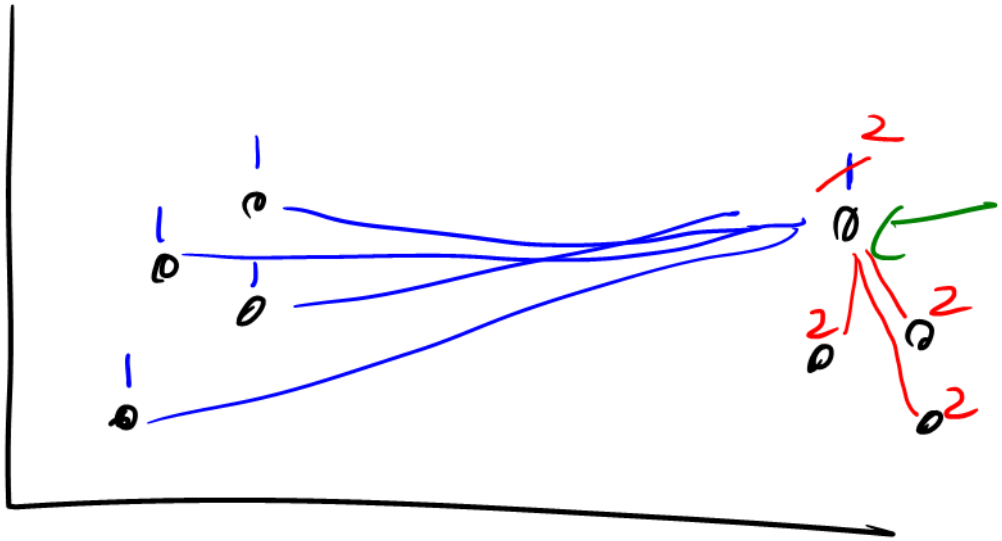

K-Means

- Lloyd's iterative refinement algorithm
  - Assign each measurement to the cluster whose mean gives the least sum of distance squared (i.e. the nearest)
  - Calculate new means to be the centroids (i.e. average) of the observations in each cluster.

See it in Python

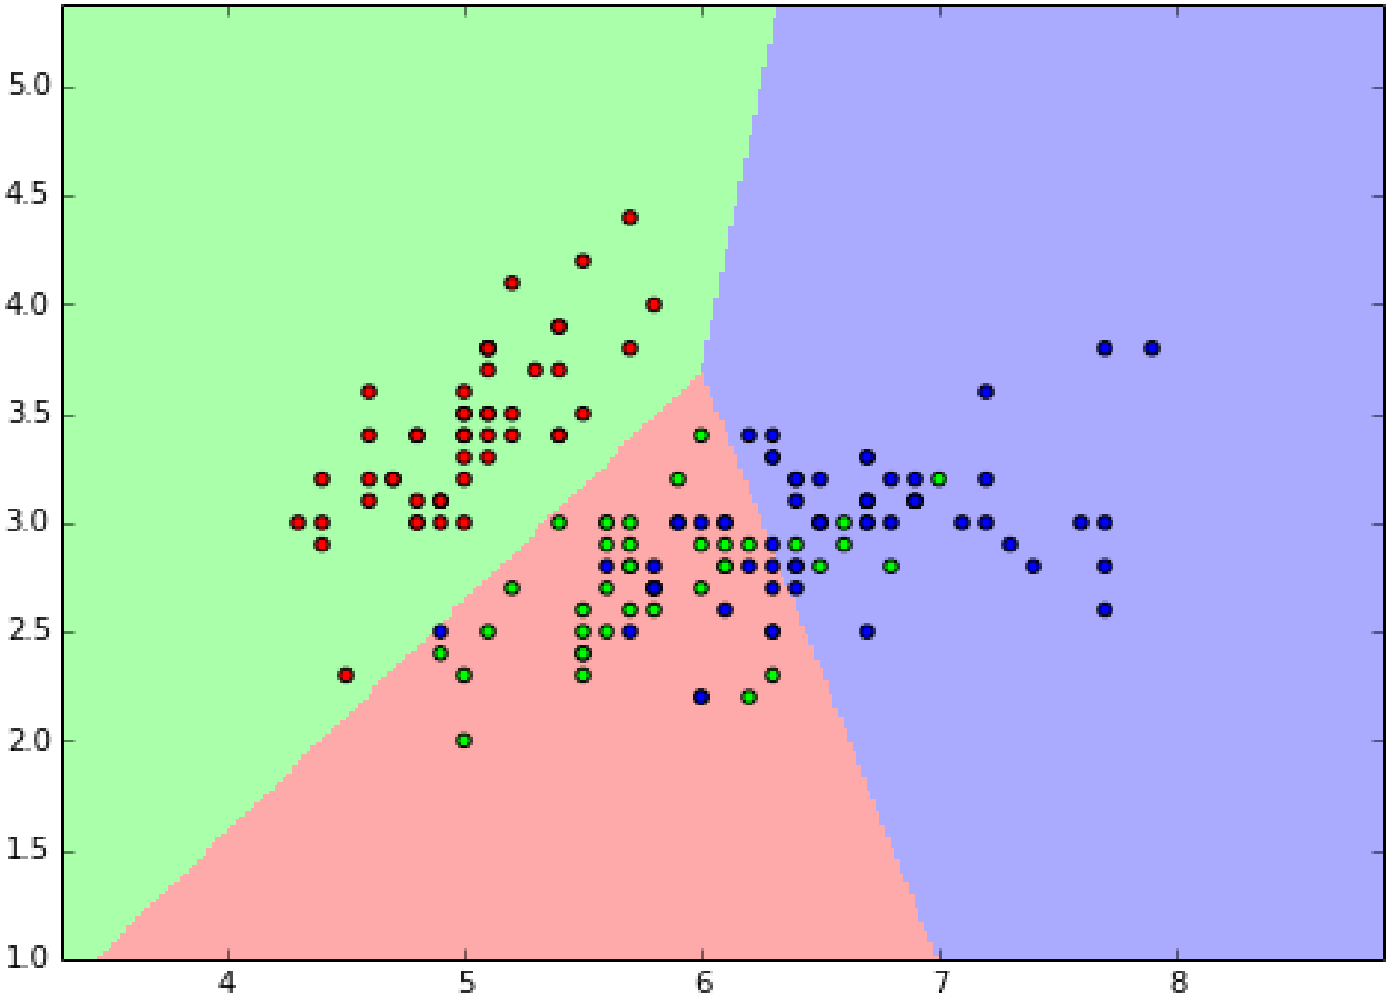

WPI

Manifold learning

- As the last item in our foray into machine learning we will dip our toes into manifold learning.
  - There are many algorithms, see Wikipedia.

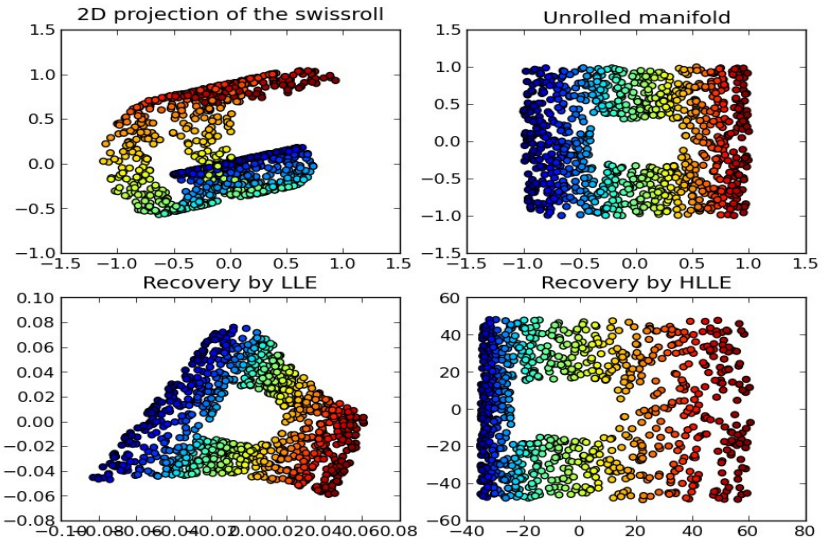

WPI

Local Tangent Space Alignment

- A very rough outline of the algorithm
  - Compute the collection of points nearest each point.
  - Compute the tangent space at each point (e.g. using PCA!)
  - Solve an optimization problem to align the tangent spaces.

See it in Python

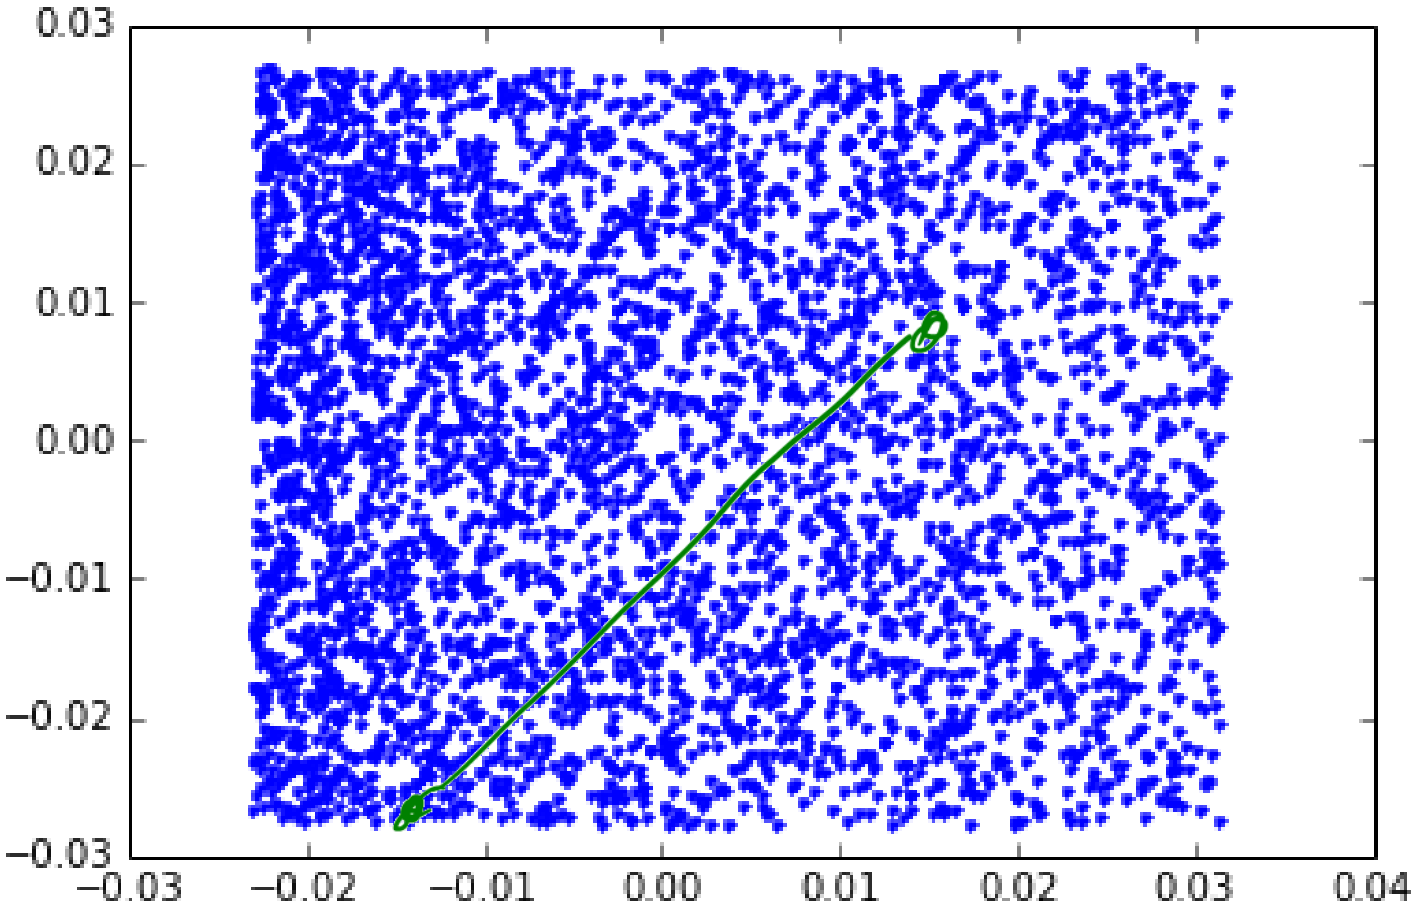

WPI